

17.14 InnoDB Startup Options and System Variables

System variables that are true or false can be enabled at server startup by naming them, or disabled by using a `--skip-` prefix. For example, to enable or disable the InnoDB adaptive hash index, you can use `--innodb-adaptive-hash-index` or `--skip-innodb-adaptive-hash-index` on the command line, or `innodb_adaptive_hash_index` or `skip_innodb_adaptive_hash_index` in an option file.

Some variable descriptions refer to “enabling” or “disabling” a variable. These variables can be enabled with the `SET` statement by setting them to `ON` or `1`, or disabled by setting them to `OFF` or `0`. Boolean variables can be set at startup to the values `ON`, `TRUE`, `OFF`, and `FALSE` (not case-sensitive), as well as `1` and `0`. See [Section 6.2.2.4, “Program Option Modifiers”](#).

System variables that take a numeric value can be specified as `--var_name=value` on the command line or as `var_name=value` in option files.

Many system variables can be changed at runtime (see [Section 7.1.9.2, “Dynamic System Variables”](#)).

For information about `GLOBAL` and `SESSION` variable scope modifiers, refer to the `SET` statement documentation.

Certain options control the locations and layout of the InnoDB data files. [Section 17.8.1, “InnoDB Startup Configuration”](#) explains how to use these options.

Some options, which you might not use initially, help tune InnoDB performance characteristics based on machine capacity and database [workload](#).

For more information on specifying options and system variables, see [Section 6.2.2, “Specifying Program Options”](#).

Table 17.23 InnoDB Option and Variable Reference

Name	Cmd-Line	Option File	System Var	Status Var	Var Scope	Dynamic
foreign_key_checks			Yes		Both	Yes
innodb_adaptive_flushing	Yes	Yes	Yes		Global	Yes
innodb_adaptive_flushing_early	Yes	Yes	Yes		Global	Yes
innodb_adaptive_hash_index	Yes	Yes	Yes		Global	Yes
innodb_adaptive_hash_index_parts	Yes	Yes	Yes		Global	No
innodb_adaptive_max_sleep_delay	Yes	Yes	Yes		Global	Yes
innodb_autoextend_increment	Yes	Yes	Yes		Global	Yes
innodb_autoinc_lock_mode	Yes	Yes	Yes		Global	No
innodb_background_drop_list_empty	Yes	Yes	Yes		Global	Yes
Innodb_buffer_pool_bytes_data				Yes	Global	No
Innodb_buffer_pool_bytes_dirty				Yes	Global	No
innodb_buffer_pool_chunk_size	Yes	Yes	Yes		Global	No
innodb_buffer_pool_debts	Yes	Yes	Yes		Global	No
innodb_buffer_pool_dump_at_shutdown	Yes	Yes	Yes		Global	Yes
innodb_buffer_pool_dump_now	Yes	Yes	Yes		Global	Yes
innodb_buffer_pool_dump_status	Yes	Yes	Yes		Global	Yes
Innodb_buffer_pool_dump_status				Yes	Global	No
innodb_buffer_pool_file_name	Yes	Yes	Yes		Global	Yes
innodb_buffer_pool_in_core_file	Yes	Yes	Yes		Global	Yes
innodb_buffer_pool_instances	Yes	Yes	Yes		Global	No
innodb_buffer_pool_load_at_startup	Yes	Yes	Yes		Global	Yes
innodb_buffer_pool_load_at_startup	Yes	Yes	Yes		Global	No
innodb_buffer_pool_load_status	Yes	Yes	Yes		Global	Yes
Innodb_buffer_pool_load_status				Yes	Global	No
Innodb_buffer_pool_pages_data				Yes	Global	No
Innodb_buffer_pool_pages_dirty				Yes	Global	No
Innodb_buffer_pool_pages_flushed				Yes	Global	No
Innodb_buffer_pool_pages_free				Yes	Global	No
Innodb_buffer_pool_pages_latched				Yes	Global	No
Innodb_buffer_pool_pages_misc				Yes	Global	No
Innodb_buffer_pool_pages_total				Yes	Global	No
Innodb_buffer_pool_read_ahead				Yes	Global	No
innodb_buffer_pool_read_ahead_evicted				Yes	Global	No
Innodb_buffer_pool_read_ahead_rnd				Yes	Global	No
Innodb_buffer_pool_read_requests				Yes	Global	No
Innodb_buffer_pool_reads				Yes	Global	No
Innodb_buffer_pool_resize_status				Yes	Global	No

innodb buffer pool size	Yes	Yes	Yes		Global	Yes
Innodb buffer pool wait free				Yes	Global	No
Innodb buffer pool write requests				Yes	Global	No
innodb change buffer max size	Yes	Yes	Yes		Global	Yes
innodb change buffering	Yes	Yes	Yes		Global	Yes
innodb change buffering debug	Yes	Yes	Yes		Global	Yes
innodb checkpoint disabled	Yes	Yes	Yes		Global	Yes
innodb checksum algorithm	Yes	Yes	Yes		Global	Yes
innodb cmp per index enabled	Yes	Yes	Yes		Global	Yes
innodb commit concurrency	Yes	Yes	Yes		Global	Yes
innodb compress debug	Yes	Yes	Yes		Global	Yes
Innodb compression failure threshold_pct	Yes	Yes	Yes		Global	Yes
innodb compression level	Yes	Yes	Yes		Global	Yes
innodb compression pad_pct_max	Yes	Yes	Yes		Global	Yes
innodb concurrency tickets	Yes	Yes	Yes		Global	Yes
innodb data file path	Yes	Yes	Yes		Global	No
Innodb data fsyncs				Yes	Global	No
innodb data home dir	Yes	Yes	Yes		Global	No
Innodb data pending fsyncs				Yes	Global	No
Innodb data pending reads				Yes	Global	No
Innodb data pending writes				Yes	Global	No
Innodb data read				Yes	Global	No
Innodb data reads				Yes	Global	No
Innodb data writes				Yes	Global	No
Innodb data written				Yes	Global	No
Innodb dblwr pages written				Yes	Global	No
Innodb dblwr writes				Yes	Global	No
innodb ddl buffer size	Yes	Yes	Yes		Session	Yes
innodb ddl log crash reset	Yes	Yes	Yes		Global	Yes
innodb ddl threads	Yes	Yes	Yes		Session	Yes
innodb deadlock detect	Yes	Yes	Yes		Global	Yes
innodb dedicated server	Yes	Yes	Yes		Global	No
innodb default row format	Yes	Yes	Yes		Global	Yes
innodb directories	Yes	Yes	Yes		Global	No
innodb disable sort file cache	Yes	Yes	Yes		Global	Yes
innodb doublewrite	Yes	Yes	Yes		Global	Yes
innodb doublewrite batch size	Yes	Yes	Yes		Global	No
innodb doublewrite dir	Yes	Yes	Yes		Global	No
innodb doublewrite files	Yes	Yes	Yes		Global	No
innodb doublewrite pages	Yes	Yes	Yes		Global	No
innodb fast shutdown	Yes	Yes	Yes		Global	Yes
innodb fil make page dirty	Yes	Yes	Yes		Global	Yes
innodb file per table	Yes	Yes	Yes		Global	Yes
innodb fill factor	Yes	Yes	Yes		Global	Yes
innodb flush log at time	Yes	Yes	Yes		Global	Yes
innodb flush log at trx commit	Yes	Yes	Yes		Global	Yes
innodb flush method	Yes	Yes	Yes		Global	No
innodb flush neighbors	Yes	Yes	Yes		Global	Yes
innodb flush sync	Yes	Yes	Yes		Global	Yes
innodb flushing avg log size	Yes	Yes	Yes		Global	Yes
innodb force load corrupt	Yes	Yes	Yes		Global	No
innodb force recovery	Yes	Yes	Yes		Global	No
innodb fsync threshold	Yes	Yes	Yes		Global	Yes
innodb ft aux table			Yes		Global	Yes
innodb ft cache size	Yes	Yes	Yes		Global	No
innodb ft enable diag plugin	Yes	Yes	Yes		Global	Yes
innodb ft enable stopwords	Yes	Yes	Yes		Both	Yes
innodb ft max token size	Yes	Yes	Yes		Global	No
innodb ft min token size	Yes	Yes	Yes		Global	No
innodb ft num word optimize	Yes	Yes	Yes		Global	Yes
innodb ft result cache limit	Yes	Yes	Yes		Global	Yes
innodb ft server stopwords	Yes	Yes	Yes		Global	Yes
innodb ft sort pll degree	Yes	Yes	Yes		Global	No
innodb ft total cache size	Yes	Yes	Yes		Global	No
innodb ft user stopwords	Yes	Yes	Yes		Both	Yes
Innodb have atomic builtins				Yes	Global	No
innodb idle flush pct	Yes	Yes	Yes		Global	Yes
innodb io capacity	Yes	Yes	Yes		Global	Yes
innodb io capacity max	Yes	Yes	Yes		Global	Yes

innodb limit optimistic insert debug	Yes	Yes	Yes		Global	Yes
innodb lock wait timeout	Yes	Yes	Yes		Both	Yes
innodb log buffer size	Yes	Yes	Yes		Global	Yes
innodb log checkpoint fuzzy	Yes	Yes	Yes		Global	Yes
innodb log checkpoint no	Yes	Yes	Yes		Global	Yes
innodb log checksums	Yes	Yes	Yes		Global	Yes
innodb log compressed pages	Yes	Yes	Yes		Global	Yes
innodb log file size	Yes	Yes	Yes		Global	No
innodb log files in group	Yes	Yes	Yes		Global	No
innodb log group home dir	Yes	Yes	Yes		Global	No
innodb log spin cpu abs	Yes	Yes	Yes		Global	Yes
innodb log spin cpu pct	Yes	Yes	Yes		Global	Yes
innodb log wait for flush spin hwm	Yes	Yes	Yes		Global	Yes
Innodb log waits				Yes	Global	No
innodb log write ahead	Yes	Yes	Yes		Global	Yes
Innodb log write requests				Yes	Global	No
innodb log writer threads	Yes	Yes	Yes		Global	Yes
Innodb log writes				Yes	Global	No
innodb lru scan depth	Yes	Yes	Yes		Global	Yes
innodb max dirty pages	Yes	Yes	Yes		Global	Yes
innodb max dirty pages pct hwm	Yes	Yes	Yes		Global	Yes
innodb max purge lag	Yes	Yes	Yes		Global	Yes
innodb max purge lag delay	Yes	Yes	Yes		Global	Yes
innodb max undo log size	Yes	Yes	Yes		Global	Yes
innodb merge threshold set a debug	Yes	Yes	Yes		Global	Yes
innodb monitor disable	Yes	Yes	Yes		Global	Yes
innodb monitor enable	Yes	Yes	Yes		Global	Yes
innodb monitor reset	Yes	Yes	Yes		Global	Yes
innodb monitor reset a debug	Yes	Yes	Yes		Global	Yes
Innodb num open files				Yes	Global	No
innodb numa interleaved	Yes	Yes	Yes		Global	No
innodb old blocks pct	Yes	Yes	Yes		Global	Yes
innodb old blocks time	Yes	Yes	Yes		Global	Yes
innodb online alter log max size	Yes	Yes	Yes		Global	Yes
innodb open files	Yes	Yes	Yes		Global	Yes
innodb optimize fulltext	Yes	Yes	Yes		Global	Yes
Innodb os log fsyncs				Yes	Global	No
Innodb os log pending fsyncs				Yes	Global	No
Innodb os log pending writes				Yes	Global	No
Innodb os log written				Yes	Global	No
innodb page cleaners	Yes	Yes	Yes		Global	No
Innodb page size				Yes	Global	No
innodb page size	Yes	Yes	Yes		Global	No
Innodb pages created				Yes	Global	No
Innodb pages read				Yes	Global	No
Innodb pages written				Yes	Global	No
innodb parallel read threads	Yes	Yes	Yes		Session	Yes
innodb print all deadlocks	Yes	Yes	Yes		Global	Yes
innodb print ddl logs	Yes	Yes	Yes		Global	Yes
innodb purge batch size	Yes	Yes	Yes		Global	Yes
innodb purge rseg truncate frequency	Yes	Yes	Yes		Global	Yes
innodb purge threads	Yes	Yes	Yes		Global	No
innodb random read ahead	Yes	Yes	Yes		Global	Yes
innodb read ahead threshold	Yes	Yes	Yes		Global	Yes
innodb read io threads	Yes	Yes	Yes		Global	No
innodb read only	Yes	Yes	Yes		Global	No
innodb redo log archive	Yes	Yes	Yes		Global	Yes
innodb redo log capacity	Yes	Yes	Yes		Global	Yes
Innodb redo log capacity resized				Yes	Global	No
Innodb redo log checkpoint lsn				Yes	Global	No
Innodb redo log current lsn				Yes	Global	No
Innodb redo log enabled				Yes	Global	No
innodb redo log encrypt	Yes	Yes	Yes		Global	Yes
innodb redo log flushed to disk lsn				Yes	Global	No
Innodb redo log logical size				Yes	Global	No
Innodb redo log physical size				Yes	Global	No
Innodb redo log read only				Yes	Global	No
Innodb redo log resize status				Yes	Global	No
Innodb redo log uuid				Yes	Global	No

innodb replication delay	Yes	Yes	Yes		Global	Yes
innodb rollback on timeout	Yes	Yes	Yes		Global	No
innodb rollback segment	Yes	Yes	Yes		Global	Yes
Innodb row lock current waits				Yes	Global	No
Innodb row lock time				Yes	Global	No
Innodb row lock time avg				Yes	Global	No
Innodb row lock time max				Yes	Global	No
Innodb row lock waits				Yes	Global	No
Innodb rows deleted				Yes	Global	No
Innodb rows inserted				Yes	Global	No
Innodb rows read				Yes	Global	No
Innodb rows updated				Yes	Global	No
innodb saved page number	Yes	Yes	Yes		Global	Yes
innodb segment reserve factor	Yes	Yes	Yes		Global	Yes
innodb sort buffer size	Yes	Yes	Yes		Global	No
innodb spin wait delay	Yes	Yes	Yes		Global	Yes
innodb spin wait pause multiplier	Yes	Yes	Yes		Global	Yes
innodb stats auto recycle	Yes	Yes	Yes		Global	Yes
innodb stats include delete marked	Yes	Yes	Yes		Global	Yes
innodb stats method	Yes	Yes	Yes		Global	Yes
innodb stats on metadata	Yes	Yes	Yes		Global	Yes
innodb stats persistent	Yes	Yes	Yes		Global	Yes
innodb stats persistent sample pages	Yes	Yes	Yes		Global	Yes
innodb stats transient sample pages	Yes	Yes	Yes		Global	Yes
innodb-status-file	Yes	Yes				
innodb status output	Yes	Yes	Yes		Global	Yes
innodb status output log	Yes	Yes	Yes		Global	Yes
innodb strict mode	Yes	Yes	Yes		Both	Yes
innodb sync array size	Yes	Yes	Yes		Global	No
innodb sync debug	Yes	Yes	Yes		Global	No
innodb sync spin loop	Yes	Yes	Yes		Global	Yes
Innodb system rows deleted				Yes	Global	No
Innodb system rows inserted				Yes	Global	No
Innodb system rows read				Yes	Global	No
Innodb system rows updated				Yes	Global	No
innodb table locks	Yes	Yes	Yes		Both	Yes
innodb temp data file path	Yes	Yes	Yes		Global	No
innodb temp tablespaces	Yes	Yes	Yes		Global	No
innodb thread concurrency	Yes	Yes	Yes		Global	Yes
innodb thread sleep delay	Yes	Yes	Yes		Global	Yes
innodb tmpdir	Yes	Yes	Yes		Both	Yes
Innodb truncated status writes				Yes	Global	No
Innodb trx purge view update delay debug	Yes	Yes	Yes		Global	Yes
innodb trx rseg n slots debug	Yes	Yes	Yes		Global	Yes
innodb undo directory	Yes	Yes	Yes		Global	No
innodb undo log encrypt	Yes	Yes	Yes		Global	Yes
innodb undo log truncate	Yes	Yes	Yes		Global	Yes
innodb undo tablespaces	Yes	Yes	Yes		Global	Yes
Innodb undo tablespaces active				Yes	Global	No
Innodb undo tablespaces explicit				Yes	Global	No
Innodb undo tablespaces implicit				Yes	Global	No
Innodb undo tablespaces total				Yes	Global	No
innodb use fdatsync	Yes	Yes	Yes		Global	Yes
innodb use native aio	Yes	Yes	Yes		Global	No
innodb validate tablespace	Yes	Yes	Yes		Global	No
innodb version			Yes		Global	No
innodb write io threads	Yes	Yes	Yes		Global	No
unique checks			Yes		Both	Yes

InnoDB Command Options

[--innodb-status-file](#)

Command-Line Format

Type

Default Value

--innodb-status-file[={OFF|ON}]

Boolean

OFF

The `--innodb-status-file` startup option controls whether InnoDB creates a file named `innodb_status.pid` in the data directory and writes [SHOW ENGINE INNODB STATUS](#) output to it every 15 seconds, approximately.

The `innodb_status.pid` file is not created by default. To create it, start [mysqld](#) with the `--innodb-status-file` option. InnoDB removes the file when the server is shut down normally. If an abnormal shutdown occurs, the status file may have to be removed manually.

The `--innodb-status-file` option is intended for temporary use, as [SHOW ENGINE INNODB STATUS](#) output generation can affect performance, and the `innodb_status.pid` file can become quite large over time.

For related information, see [Section 17.17.2, “Enabling InnoDB Monitors”](#).

InnoDB System Variables

[innodb_adaptive_flushing](#)

Command-Line Format	<code>--innodb-adaptive-flushing[={OFF ON}]</code>
System Variable	innodb_adaptive_flushing
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	ON

Specifies whether to dynamically adjust the rate of flushing [dirty pages](#) in the InnoDB [buffer pool](#) based on the workload. Adjusting the flush rate dynamically is intended to avoid bursts of I/O activity. This setting is enabled by default. See [Section 17.8.3.5, “Configuring Buffer Pool Flushing”](#) for more information. For general I/O tuning advice, see [Section 10.5.8, “Optimizing InnoDB Disk I/O”](#).

[innodb_adaptive_flushing_lwm](#)

Command-Line Format	<code>--innodb-adaptive-flushing-lwm=#</code>
System Variable	innodb_adaptive_flushing_lwm
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	10
Minimum Value	0
Maximum Value	70

Defines the low water mark representing percentage of [redo log](#) capacity at which [adaptive flushing](#) is enabled. For more information, see [Section 17.8.3.5, “Configuring Buffer Pool Flushing”](#).

[innodb_adaptive_hash_index](#)

Command-Line Format	<code>--innodb-adaptive-hash-index[={OFF ON}]</code>
System Variable	innodb_adaptive_hash_index
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	OFF

Whether the InnoDB [adaptive hash index](#) is enabled or disabled. It may be desirable, depending on your workload, to dynamically enable or disable [adaptive hash indexing](#) to improve query performance. Because the adaptive hash index may not be useful for all workloads, conduct benchmarks with it both enabled and disabled, using realistic workloads. See [Section 17.5.3, “Adaptive Hash Index”](#) for details.

This variable is disabled by default. You can modify this parameter using the `SET GLOBAL` statement, without restarting the server. Changing the setting at runtime requires privileges sufficient to set global system variables. See [Section 7.1.9.1, “System Variable Privileges”](#). You can also use [--innodb-adaptive-hash-index](#) at server startup to enable it.

Disabling the adaptive hash index empties the hash table immediately. Normal operations can continue while the hash table is emptied, and executing queries that were using the hash table access the index B-trees directly instead. When the adaptive hash index is re-enabled, the hash table is populated again during normal operation.

Before MySQL 8.4, this option was enabled by default.

[innodb_adaptive_hash_index_parts](#)

Command-Line Format	<code>--innodb-adaptive-hash-index-parts=#</code>
System Variable	innodb_adaptive_hash_index_parts
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	Numeric
Default Value	8

Minimum Value	1
Maximum Value	512

Partitions the adaptive hash index search system. Each index is bound to a specific partition, with each partition protected by a separate latch.

The adaptive hash index search system is partitioned into 8 parts by default. The maximum setting is 512.

For related information, see [Section 17.5.3, “Adaptive Hash Index”](#).

[innodb_adaptive_max_sleep_delay](#)

Command-Line Format	--innodb-adaptive-max-sleep-delay=#
System Variable	innodb_adaptive_max_sleep_delay
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	150000
Minimum Value	0
Maximum Value	1000000
Unit	microseconds

Permits InnoDB to automatically adjust the value of [innodb_thread_sleep_delay](#) up or down according to the current workload. Any nonzero value enables automated, dynamic adjustment of the [innodb_thread_sleep_delay](#) value, up to the maximum value specified in the [innodb_adaptive_max_sleep_delay](#) option. The value represents the number of microseconds. This option can be useful in busy systems, with greater than 16 InnoDB threads. (In practice, it is most valuable for MySQL systems with hundreds or thousands of simultaneous connections.)

For more information, see [Section 17.8.4, “Configuring Thread Concurrency for InnoDB”](#).

[innodb_autoextend_increment](#)

Command-Line Format	--innodb-autoextend-increment=#
System Variable	innodb_autoextend_increment
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	64
Minimum Value	1
Maximum Value	1000
Unit	megabytes

The increment size (in megabytes) for extending the size of an auto-extending InnoDB [system tablespace](#) file when it becomes full. The default value is 64. For related information, see [System Tablespace Data File Configuration](#), and [Resizing the System Tablespace](#).

The [innodb_autoextend_increment](#) setting does not affect [file-per-table](#) tablespace files or [general tablespace](#) files. These files are auto-extending regardless of the [innodb_autoextend_increment](#) setting. The initial extensions are by small amounts, after which extensions occur in increments of 4MB.

[innodb_autoinc_lock_mode](#)

Command-Line Format	--innodb-autoinc-lock-mode=#
System Variable	innodb_autoinc_lock_mode
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	Integer
Default Value	2
	0
Valid Values	1
	2

The [lock mode](#) to use for generating [auto-increment](#) values. Permissible values are 0, 1, or 2, for traditional, consecutive, or interleaved, respectively.

The default setting is 2 (interleaved), for compatibility with row-based replication.

For the characteristics of each lock mode, see [InnoDB AUTO INCREMENT Lock Modes](#).

[innodb_background_drop_list_empty](#)

Command-Line Format	--innodb-background-drop-list-empty[={OFF ON}]
----------------------------	--

System Variable	<u>innodb_background_drop_list_empty</u>
Scope	Global
Dynamic	Yes
<u>SET VAR</u> Hint Applies	No
Type	Boolean
Default Value	OFF

Enabling the [innodb_background_drop_list_empty](#) debug option helps avoid test case failures by delaying table creation until the background drop list is empty. For example, if test case A places table t1 on the background drop list, test case B waits until the background drop list is empty before creating table t1.

[innodb_buffer_pool_chunk_size](#)

Command-Line Format	--innodb-buffer-pool-chunk-size=#
System Variable	<u>innodb_buffer_pool_chunk_size</u>
Scope	Global
Dynamic	No
<u>SET VAR</u> Hint Applies	No
Type	Integer
Default Value	134217728
Minimum Value	1048576
Maximum Value	innodb_buffer_pool_size / innodb_buffer_pool_instances
Unit	bytes

[innodb_buffer_pool_chunk_size](#) defines the chunk size for InnoDB buffer pool resizing operations.

To avoid copying all buffer pool pages during resizing operations, the operation is performed in “chunks”. By default, [innodb_buffer_pool_chunk_size](#) is 128MB (134217728 bytes). The number of pages contained in a chunk depends on the value of [innodb_page_size](#).

[innodb_buffer_pool_chunk_size](#) can be increased or decreased in units of 1MB (1048576 bytes).

The following conditions apply when altering the [innodb_buffer_pool_chunk_size](#) value:

If $\text{innodb_buffer_pool_chunk_size} * \text{innodb_buffer_pool_instances}$ is larger than the current buffer pool size when the buffer pool is initialized, [innodb_buffer_pool_chunk_size](#) is truncated to $\text{innodb_buffer_pool_size} / \text{innodb_buffer_pool_instances}$.

Buffer pool size must always be equal to or a multiple of $\text{innodb_buffer_pool_chunk_size} * \text{innodb_buffer_pool_instances}$. If you alter [innodb_buffer_pool_chunk_size](#), [innodb_buffer_pool_size](#) is automatically rounded to a value that is equal to or a multiple of $\text{innodb_buffer_pool_chunk_size} * \text{innodb_buffer_pool_instances}$. The adjustment occurs when the buffer pool is initialized.

Important

Care should be taken when changing [innodb_buffer_pool_chunk_size](#), as changing this value can automatically increase the size of the buffer pool. Before changing [innodb_buffer_pool_chunk_size](#), calculate its effect on [innodb_buffer_pool_size](#) to ensure that the resulting buffer pool size is acceptable.

To avoid potential performance issues, the number of chunks ($\text{innodb_buffer_pool_size} / \text{innodb_buffer_pool_chunk_size}$) should not exceed 1000.

The [innodb_buffer_pool_size](#) variable is dynamic, which permits resizing the buffer pool while the server is online. However, the buffer pool size must be equal to or a multiple of $\text{innodb_buffer_pool_chunk_size} * \text{innodb_buffer_pool_instances}$, and changing either of those variable settings requires restarting the server.

See [Section 17.8.3.1, “Configuring InnoDB Buffer Pool Size”](#) for more information.

[innodb_buffer_pool_debug](#)

Command-Line Format	--innodb-buffer-pool-debug[={OFF ON}]
System Variable	<u>innodb_buffer_pool_debug</u>
Scope	Global
Dynamic	No
<u>SET VAR</u> Hint Applies	No
Type	Boolean
Default Value	OFF

Enabling this option permits multiple buffer pool instances when the buffer pool is less than 1GB in size, ignoring the 1GB minimum buffer pool size constraint imposed on [innodb_buffer_pool_instances](#). The [innodb_buffer_pool_debug](#) option is only available if debugging support is compiled in using the [WITH_DEBUG](#) CMake option.

[innodb_buffer_pool_dump_at_shutdown](#)

Command-Line Format	--innodb-buffer-pool-dump-at-shutdown[={OFF ON}]
System Variable	<u>innodb_buffer_pool_dump_at_shutdown</u>
Scope	Global

Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	ON

Specifies whether to record the pages cached in the InnoDB [buffer pool](#) when the MySQL server is shut down, to shorten the [warmup](#) process at the next restart. Typically used in combination with [innodb_buffer_pool_load_at_startup](#). The [innodb_buffer_pool_dump_pct](#) option defines the percentage of most recently used buffer pool pages to dump.

Both [innodb_buffer_pool_dump_at_shutdown](#) and [innodb_buffer_pool_load_at_startup](#) are enabled by default.

For more information, see [Section 17.8.3.6, “Saving and Restoring the Buffer Pool State”](#).

[innodb_buffer_pool_dump_now](#)

Command-Line Format	--innodb-buffer-pool-dump-now[={OFF ON}]
System Variable	innodb_buffer_pool_dump_now
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	OFF

Immediately makes a record of pages cached in the InnoDB [buffer pool](#). Typically used in combination with [innodb_buffer_pool_load_now](#).

Enabling [innodb_buffer_pool_dump_now](#) triggers the recording action but does not alter the variable setting, which always remains OFF or 0. To view buffer pool dump status after triggering a dump, query the [InnoDB_buffer_pool_dump_status](#) variable.

Enabling [innodb_buffer_pool_dump_now](#) triggers the dump action but does not alter the variable setting, which always remains OFF or 0. To view buffer pool dump status after triggering a dump, query the [InnoDB_buffer_pool_dump_status](#) variable.

For more information, see [Section 17.8.3.6, “Saving and Restoring the Buffer Pool State”](#).

[innodb_buffer_pool_dump_pct](#)

Command-Line Format	--innodb-buffer-pool-dump-pct=#
System Variable	innodb_buffer_pool_dump_pct
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	25
Minimum Value	1
Maximum Value	100

Specifies the percentage of the most recently used pages for each buffer pool to read out and dump. The range is 1 to 100. The default value is 25. For example, if there are 4 buffer pools with 100 pages each, and [innodb_buffer_pool_dump_pct](#) is set to 25, the 25 most recently used pages from each buffer pool are dumped.

[innodb_buffer_pool_filename](#)

Command-Line Format	--innodb-buffer-pool-filename=file_name
System Variable	innodb_buffer_pool_filename
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	File name
Default Value	ib_buffer_pool

Specifies the name of the file that holds the list of tablespace IDs and page IDs produced by [innodb_buffer_pool_dump_at_shutdown](#) or [innodb_buffer_pool_dump_now](#). Tablespace IDs and page IDs are saved in the following format: space, page_id. By default, the file is named ib_buffer_pool and is located in the InnoDB data directory. A non-default location must be specified relative to the data directory.

A file name can be specified at runtime, using a [SET](#) statement:

```
SET GLOBAL innodb_buffer_pool_filename='file_name';
```

You can also specify a file name at startup, in a startup string or MySQL configuration file. When specifying a file name at startup, the file must exist or InnoDB returns a startup error indicating that there is no such file or directory.

For more information, see [Section 17.8.3.6, “Saving and Restoring the Buffer Pool State”](#).

[innodb_buffer_pool_in_core_file](#)

Command-Line Format	--innodb-buffer-pool-in-core-file[={OFF ON}]
System Variable	innodb_buffer_pool_in_core_file
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	OFF

Disabling (default) the [innodb_buffer_pool_in_core_file](#) variable reduces the size of core files by excluding InnoDB buffer pool pages.

To use this variable, the [core_file](#) variable must be enabled, and to disable this option the operating system must support the MADV_DONTDUMP non-POSIX extension to `madvise()`, which is supported in Linux 3.4 and later. For more information, see [Section 17.8.3.7, “Excluding or Including Buffer Pool Pages from Core Files”](#).

This is disabled by default on systems that support MADV_DONTDUMP, which is typically only Linux and not macOS or Windows.

Before MySQL 8.4, this option was enabled by default.

[innodb_buffer_pool_instances](#)

Command-Line Format	--innodb-buffer-pool-instances=#
System Variable	innodb_buffer_pool_instances
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	Integer
Default Value	see description
Minimum Value	1
Maximum Value	64

The number of regions that the InnoDB [buffer pool](#) is divided into. For systems with buffer pools in the multi-gigabyte range, dividing the buffer pool into separate instances can improve concurrency, by reducing contention as different threads read and write to cached pages. Each page that is stored in or read from the buffer pool is assigned to one of the buffer pool instances randomly, using a hashing function. Each buffer pool manages its own free lists, [flush lists](#), [LRUs](#), and all other data structures connected to a buffer pool, and is protected by its own buffer pool [mutex](#).

The total buffer pool size is divided among all the buffer pools. For best efficiency, specify a combination of [innodb_buffer_pool_instances](#) and [innodb_buffer_pool_size](#) so that each buffer pool instance is at least 1GB.

If [innodb_buffer_pool_size](#) ≤ 1 GiB, then the default [innodb_buffer_pool_instances](#) value is 1.

If [innodb_buffer_pool_size](#) > 1 GiB, then the default [innodb_buffer_pool_instances](#) value is the minimum value from the following two calculated hints, within a range of 1-64:

Buffer pool hint: calculated as 1/2 of ([innodb_buffer_pool_size](#) / [innodb_buffer_pool_chunk_size](#))

CPU hint: calculated as 1/4 of available logical processors

For related information, see [Section 17.8.3.1, “Configuring InnoDB Buffer Pool Size”](#).

[innodb_buffer_pool_load_abort](#)

Command-Line Format	--innodb-buffer-pool-load-abort[={OFF ON}]
System Variable	innodb_buffer_pool_load_abort
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	OFF

Interrupts the process of restoring InnoDB [buffer pool](#) contents triggered by [innodb_buffer_pool_load_at_startup](#) or [innodb_buffer_pool_load_now](#).

Enabling [innodb_buffer_pool_load_abort](#) triggers the abort action but does not alter the variable setting, which always remains OFF or 0. To view buffer pool load status after triggering an abort action, query the [innodb_buffer_pool_load_status](#) variable.

For more information, see [Section 17.8.3.6, “Saving and Restoring the Buffer Pool State”](#).

[innodb_buffer_pool_load_at_startup](#)

Command-Line Format	--innodb-buffer-pool-load-at-startup[={OFF ON}]
System Variable	innodb_buffer_pool_load_at_startup
Scope	Global

Dynamic	No
SET VAR Hint Applies	No
Type	Boolean
Default Value	ON

Specifies that, on MySQL server startup, the InnoDB [buffer pool](#) is automatically [warmed up](#) by loading the same pages it held at an earlier time. Typically used in combination with [innodb_buffer_pool_dump_at_shutdown](#).

Both [innodb_buffer_pool_dump_at_shutdown](#) and [innodb_buffer_pool_load_at_startup](#) are enabled by default.

For more information, see [Section 17.8.3.6, “Saving and Restoring the Buffer Pool State”](#).

[innodb_buffer_pool_load_now](#)

Command-Line Format	--innodb-buffer-pool-load-now[={OFF ON}]
System Variable	innodb_buffer_pool_load_now
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	OFF

Immediately [warms up](#) the InnoDB [buffer pool](#) by loading data pages without waiting for a server restart. Can be useful to bring cache memory back to a known state during benchmarking or to ready the MySQL server to resume its normal workload after running queries for reports or maintenance.

Enabling [innodb_buffer_pool_load_now](#) triggers the load action but does not alter the variable setting, which always remains OFF or 0. To view buffer pool load progress after triggering a load, query the [innodb_buffer_pool_load_status](#) variable.

For more information, see [Section 17.8.3.6, “Saving and Restoring the Buffer Pool State”](#).

[innodb_buffer_pool_size](#)

Command-Line Format	--innodb-buffer-pool-size=#
System Variable	innodb_buffer_pool_size
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	134217728
Minimum Value	5242880
Maximum Value (64-bit platforms)	2**64-1
Maximum Value (32-bit platforms)	2**32-1
Unit	bytes

The size in bytes of the [buffer pool](#), the memory area where InnoDB caches table and index data. The default value is 134217728 bytes (128MB). The maximum value depends on the CPU architecture; the maximum is 4294967295 ($2^{32}-1$) on 32-bit systems and 18446744073709551615 ($2^{64}-1$) on 64-bit systems. On 32-bit systems, the CPU architecture and operating system may impose a lower practical maximum size than the stated maximum. When the size of the buffer pool is greater than 1GB, setting [innodb_buffer_pool_instances](#) to a value greater than 1 can improve the scalability on a busy server.

A larger buffer pool requires less disk I/O to access the same table data more than once. On a dedicated database server, you might set the buffer pool size to 80% of the machine's physical memory size. Be aware of the following potential issues when configuring buffer pool size, and be prepared to scale back the size of the buffer pool if necessary.

Competition for physical memory can cause paging in the operating system.

InnoDB reserves additional memory for buffers and control structures, so that the total allocated space is approximately 10% greater than the specified buffer pool size.

Address space for the buffer pool must be contiguous, which can be an issue on Windows systems with DLLs that load at specific addresses.

The time to initialize the buffer pool is roughly proportional to its size. On instances with large buffer pools, initialization time might be significant. To reduce the initialization period, you can save the buffer pool state at server shutdown and restore it at server startup. See [Section 17.8.3.6, “Saving and Restoring the Buffer Pool State”](#).

When you increase or decrease buffer pool size, the operation is performed in chunks. Chunk size is defined by the [innodb_buffer_pool_chunk_size](#) variable, which has a default of 128 MB.

Buffer pool size must always be equal to or a multiple of [innodb_buffer_pool_chunk_size](#) * [innodb_buffer_pool_instances](#). If you alter the buffer pool size to a value that is not equal to or a multiple of [innodb_buffer_pool_chunk_size](#) * [innodb_buffer_pool_instances](#), buffer pool size is automatically adjusted to a value that is equal to or a multiple of [innodb_buffer_pool_chunk_size](#) * [innodb_buffer_pool_instances](#).

[innodb_buffer_pool_size](#) can be set dynamically, which allows you to resize the buffer pool without restarting the server. The [Innodb_buffer_pool_resize_status](#) status variable reports the status of online buffer pool resizing operations. See [Section 17.8.3.1, “Configuring InnoDB Buffer Pool Size”](#) for more information.

If [innodb_dedicated_server](#) is enabled, the [innodb_buffer_pool_size](#) value is automatically configured if it is not explicitly defined. For more information, see [Section 17.8.12, “Enabling Automatic Configuration for a Dedicated MySQL Server”](#).

[innodb_change_buffer_max_size](#)

Command-Line Format	--innodb-change-buffer-max-size=#
System Variable	innodb_change_buffer_max_size
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	25
Minimum Value	0
Maximum Value	50

Maximum size for the InnoDB [change buffer](#), as a percentage of the total size of the [buffer pool](#). You might increase this value for a MySQL server with heavy insert, update, and delete activity, or decrease it for a MySQL server with unchanging data used for reporting. For more information, see [Section 17.5.2, “Change Buffer”](#). For general I/O tuning advice, see [Section 10.5.8, “Optimizing InnoDB Disk I/O”](#).

[innodb_change_buffering](#)

Command-Line Format	--innodb-change-buffering=value
System Variable	innodb_change_buffering
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Enumeration
Default Value	none
	none
	inserts
	deletes
Valid Values	changes
	purges
	all

Whether InnoDB performs [change buffering](#), an optimization that delays write operations to secondary indexes so that the I/O operations can be performed sequentially. Permitted values are described in the following table. Values may also be specified numerically.

Table 17.24 Permitted Values for innodb_change_buffering

Value		Numeric Value	Description
none	0		Default. Do not buffer any operations.
inserts	1		Buffer insert operations.
deletes	2		Buffer delete marking operations; strictly speaking, the writes that mark index records for later deletion during a purge operation.
changes	3		Buffer inserts and delete-marking operations.
purges	4		Buffer the physical deletion operations that happen in the background.
all	5		Buffer inserts, delete-marking operations, and purges.

Before MySQL 8.4, the default value was `all`.

For more information, see [Section 17.5.2, “Change Buffer”](#). For general I/O tuning advice, see [Section 10.5.8, “Optimizing InnoDB Disk I/O”](#).

[innodb_change_buffering_debug](#)

Command-Line Format	--innodb-change-buffering-debug=#
System Variable	innodb_change_buffering_debug
Scope	Global
Dynamic	Yes

SET VAR Hint Applies	No
Type	Integer
Default Value	0
Minimum Value	0
Maximum Value	2

Sets a debug flag for InnoDB change buffering. A value of 1 forces all changes to the change buffer. A value of 2 causes an unexpected exit at merge. A default value of 0 indicates that the change buffering debug flag is not set. This option is only available when debugging support is compiled in using the [WITH_DEBUG](#) CMake option.

[innodb_checkpoint_disabled](#)

Command-Line Format	--innodb-checkpoint-disabled[={OFF ON}]
System Variable	innodb_checkpoint_disabled
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	OFF

This is a debug option that is only intended for expert debugging use. It disables checkpoints so that a deliberate server exit always initiates InnoDB recovery. It should only be enabled for a short interval, typically before running DML operations that write redo log entries that would require recovery following a server exit. This option is only available if debugging support is compiled in using the [WITH_DEBUG](#) CMake option.

[innodb_checksum_algorithm](#)

Command-Line Format	--innodb-checksum-algorithm=value
System Variable	innodb_checksum_algorithm
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Enumeration
Default Value	crc32
	crc32
	strict_crc32
	innodb
Valid Values	strict_innodb
	none
	strict_none

Specifies how to generate and verify the [checksum](#) stored in the disk blocks of InnoDB [tablespaces](#). The default value for [innodb_checksum_algorithm](#) is `crc32`.

The value `innodb` is backward-compatible with earlier versions of MySQL. The value `crc32` uses an algorithm that is faster to compute the checksum for every modified block, and to check the checksums for each disk read. It scans blocks 64 bits at a time, which is faster than the `innodb` checksum algorithm, which scans blocks 8 bits at a time. The value `none` writes a constant value in the checksum field rather than computing a value based on the block data. The blocks in a tablespace can use a mix of old, new, and no checksum values, being updated gradually as the data is modified; once blocks in a tablespace are modified to use the `crc32` algorithm, the associated tables cannot be read by earlier versions of MySQL.

The strict form of a checksum algorithm reports an error if it encounters a valid but non-matching checksum value in a tablespace. It is recommended that you only use strict settings in a new instance, to set up tablespaces for the first time. Strict settings are somewhat faster, because they do not need to compute all checksum values during disk reads.

The following table shows the difference between the `none`, `innodb`, and `crc32` option values, and their strict counterparts. `none`, `innodb`, and `crc32` write the specified type of checksum value into each data block, but for compatibility accept other checksum values when verifying a block during a read operation. Strict settings also accept valid checksum values but print an error message when a valid non-matching checksum value is encountered. Using the strict form can make verification faster if all InnoDB data files in an instance are created under an identical [innodb_checksum_algorithm](#) value.

Table 17.25 Permitted innodb_checksum_algorithm Values

Value	Generated checksum (when writing)	Permitted checksums (when reading)
<code>none</code>	A constant number.	Any of the checksums generated by <code>none</code> , <code>innodb</code> , or <code>crc32</code> .
<code>innodb</code>	A checksum calculated in software, using the original algorithm from InnoDB.	Any of the checksums generated by <code>none</code> , <code>innodb</code> , or <code>crc32</code> .
<code>crc32</code>	A checksum calculated using the <code>crc32</code> algorithm, possibly done with a hardware assist.	Any of the checksums generated by <code>none</code> , <code>innodb</code> , or <code>crc32</code> .

strict_none	A constant number	Any of the checksums generated by <code>none</code> , <code>innodb</code> , or <code>crc32</code> . InnoDB prints an error message if a valid but non-matching checksum is encountered.
strict_innodb	A checksum calculated in software, using the original algorithm from InnoDB.	Any of the checksums generated by <code>none</code> , <code>innodb</code> , or <code>crc32</code> . InnoDB prints an error message if a valid but non-matching checksum is encountered.
strict_crc32	A checksum calculated using the <code>crc32</code> algorithm, possibly done with a hardware assist.	Any of the checksums generated by <code>none</code> , <code>innodb</code> , or <code>crc32</code> . InnoDB prints an error message if a valid but non-matching checksum is encountered.

[innodb_cmp_per_index_enabled](#)

Command-Line Format	<code>--innodb-cmp-per-index-enabled[={OFF ON}]</code>
System Variable	innodb_cmp_per_index_enabled
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	OFF

Enables per-index compression-related statistics in the Information Schema [INNODB_CMP_PER_INDEX](#) table. Because these statistics can be expensive to gather, only enable this option on development, test, or replica instances during performance tuning related to InnoDB [compressed](#) tables.

For more information, see [Section 28.4.8, “The INFORMATION_SCHEMA INNODB_CMP_PER_INDEX and INNODB_CMP_PER_INDEX_RESET Tables”](#), and [Section 17.9.1.4, “Monitoring InnoDB Table Compression at Runtime”](#).

[innodb_commit_concurrency](#)

Command-Line Format	<code>--innodb-commit-concurrency=#</code>
System Variable	innodb_commit_concurrency
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	0
Minimum Value	0
Maximum Value	1000

The number of [threads](#) that can [commit](#) at the same time. A value of 0 (the default) permits any number of [transactions](#) to commit simultaneously.

The value of [innodb_commit_concurrency](#) cannot be changed at runtime from zero to nonzero or vice versa. The value can be changed from one nonzero value to another.

[innodb_compress_debug](#)

Command-Line Format	<code>--innodb-compress-debug=value</code>
System Variable	innodb_compress_debug
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Enumeration
Default Value	<code>none</code>
	<code>none</code>
Valid Values	<code>zlib</code>
	<code>lz4</code>
	<code>lz4hc</code>

Compresses all tables using a specified compression algorithm without having to define a `COMPRESSION` attribute for each table. This option is only available if debugging support is compiled in using the [WITH_DEBUG CMake](#) option.

For related information, see [Section 17.9.2, “InnoDB Page Compression”](#).

[innodb_compression_failure_threshold_pct](#)

Command-Line Format	<code>--innodb-compression-failure-threshold-pct=#</code>
----------------------------	---

System Variable	<u>innodb_compression_failure_threshold_pct</u>
Scope	Global
Dynamic	Yes
<u>SET VAR</u> Hint Applies	No
Type	Integer
Default Value	5
Minimum Value	0
Maximum Value	100

Defines the compression failure rate threshold for a table, as a percentage, at which point MySQL begins adding padding within [compressed](#) pages to avoid expensive [compression failures](#). When this threshold is passed, MySQL begins to leave additional free space within each new compressed page, dynamically adjusting the amount of free space up to the percentage of page size specified by [innodb_compression_pad_pct_max](#). A value of zero disables the mechanism that monitors compression efficiency and dynamically adjusts the padding amount.

For more information, see [Section 17.9.1.6, “Compression for OLTP Workloads”](#).

[innodb_compression_level](#)

Command-Line Format	--innodb-compression-level=#
System Variable	<u>innodb_compression_level</u>
Scope	Global
Dynamic	Yes
<u>SET VAR</u> Hint Applies	No
Type	Integer
Default Value	6
Minimum Value	0
Maximum Value	9

Specifies the level of zlib compression to use for InnoDB [compressed](#) tables and indexes. A higher value lets you fit more data onto a storage device, at the expense of more CPU overhead during compression. A lower value lets you reduce CPU overhead when storage space is not critical, or you expect the data is not especially compressible.

For more information, see [Section 17.9.1.6, “Compression for OLTP Workloads”](#).

[innodb_compression_pad_pct_max](#)

Command-Line Format	--innodb-compression-pad-pct-max=#
System Variable	<u>innodb_compression_pad_pct_max</u>
Scope	Global
Dynamic	Yes
<u>SET VAR</u> Hint Applies	No
Type	Integer
Default Value	50
Minimum Value	0
Maximum Value	75

Specifies the maximum percentage that can be reserved as free space within each compressed [page](#), allowing room to reorganize the data and modification log within the page when a [compressed](#) table or index is updated and the data might be recompressed. Only applies when [innodb_compression_failure_threshold_pct](#) is set to a nonzero value, and the rate of [compression failures](#) passes the cutoff point.

For more information, see [Section 17.9.1.6, “Compression for OLTP Workloads”](#).

[innodb_concurrency_tickets](#)

Command-Line Format	--innodb-concurrency-tickets=#
System Variable	<u>innodb_concurrency_tickets</u>
Scope	Global
Dynamic	Yes
<u>SET VAR</u> Hint Applies	No
Type	Integer
Default Value	5000
Minimum Value	1
Maximum Value	4294967295

Determines the number of [threads](#) that can enter InnoDB concurrently. A thread is placed in a queue when it tries to enter InnoDB if the number of threads has already reached the concurrency limit. When a thread is permitted to enter InnoDB, it is given a number of “tickets” equal to the value of [innodb_concurrency_tickets](#), and the thread can enter and leave InnoDB freely until it has used up its tickets. After that point, the thread again becomes subject to the concurrency check (and possible queuing) the next time it tries to enter InnoDB. The default value is 5000.

With a small [innodb_concurrency_tickets](#) value, small transactions that only need to process a few rows compete fairly with larger transactions that process many rows. The disadvantage of a small [innodb_concurrency_tickets](#) value is that large transactions must loop through the queue many times before they can complete, which extends the amount of time required to complete their task.

With a large [innodb_concurrency_tickets](#) value, large transactions spend less time waiting for a position at the end of the queue (controlled by [innodb_thread_concurrency](#)) and more time retrieving rows. Large transactions also require fewer trips through the queue to complete their task. The disadvantage of a large [innodb_concurrency_tickets](#) value is that too many large transactions running at the same time can starve smaller transactions by making them wait a longer time before executing.

With a nonzero [innodb_thread_concurrency](#) value, you may need to adjust the [innodb_concurrency_tickets](#) value up or down to find the optimal balance between larger and smaller transactions. The `SHOW ENGINE INNODB STATUS` report shows the number of tickets remaining for an executing transaction in its current pass through the queue. This data may also be obtained from the `TRX_CONCURRENCY_TICKETS` column of the Information Schema [INNODB_TRX](#) table.

For more information, see [Section 17.8.4, “Configuring Thread Concurrency for InnoDB”](#).

[innodb_data_file_path](#)

Command-Line Format	<code>--innodb-data-file-path=file_name</code>
System Variable	innodb_data_file_path
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	String
Default Value	<code>ibdata1:12M:autoextend</code>

Defines the name, size, and attributes of InnoDB system tablespace data files. If you do not specify a value for [innodb_data_file_path](#), the default behavior is to create a single auto-extending data file, slightly larger than 12MB, named `ibdata1`.

The full syntax for a data file specification includes the file name, file size, `autoextend` attribute, and `max` attribute:

```
file_name:file_size[:autoextend[:max:max_file_size]]
```

File sizes are specified in kilobytes, megabytes, or gigabytes by appending `K`, `M` or `G` to the size value. If specifying the data file size in kilobytes, do so in multiples of 1024. Otherwise, KB values are rounded to nearest megabyte (MB) boundary. The sum of file sizes must be, at a minimum, slightly larger than 12MB.

For additional configuration information, see [System Tablespace Data File Configuration](#). For resizing instructions, see [Resizing the System Tablespace](#).

[innodb_data_home_dir](#)

Command-Line Format	<code>--innodb-data-home-dir=dir_name</code>
System Variable	innodb_data_home_dir
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	Directory name

The common part of the directory path for InnoDB [system tablespace](#) data files. The default value is the MySQL data directory. The setting is concatenated with the [innodb_data_file_path](#) setting, unless that setting is defined with an absolute path.

A trailing slash is required when specifying a value for [innodb_data_home_dir](#). For example:

```
[mysqld]
innodb_data_home_dir = /path/to/myibdata/
```

This setting does not affect the location of [file-per-table](#) tablespaces.

For related information, see [Section 17.8.1, “InnoDB Startup Configuration”](#).

[innodb_ddl_buffer_size](#)

Command-Line Format	<code>--innodb-ddl-buffer-size=#</code>
System Variable	innodb_ddl_buffer_size
Scope	Session
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	1048576
Minimum Value	65536
Maximum Value	4294967295
Unit	bytes

Defines the maximum buffer size for DDL operations. The default setting is 1048576 bytes (approximately 1 MB). Applies to online DDL operations that create or rebuild secondary indexes. See [Section 17.12.4, “Online DDL Memory Management”](#). The maximum buffer size per DDL thread is the maximum buffer size divided by the number of DDL threads ([innodb_ddl_buffer_size/innodb_ddl_threads](#)).

[innodb_ddl_log_crash_reset_debug](#)

Command-Line Format	--innodb-ddl-log-crash-reset-debug[={OFF ON}]
System Variable	innodb_ddl_log_crash_reset_debug
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	OFF

Enable this debug option to reset DDL log crash injection counters to 1. This option is only available when debugging support is compiled in using the [WITH_DEBUG](#) CMake option.

[innodb_ddl_threads](#)

Command-Line Format	--innodb-ddl-threads=#
System Variable	innodb_ddl_threads
Scope	Session
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	4
Minimum Value	1
Maximum Value	64

Defines the maximum number of parallel threads for the sort and build phases of index creation. Applies to online DDL operations that create or rebuild secondary indexes. For related information, see [Section 17.12.5, “Configuring Parallel Threads for Online DDL Operations”](#), and [Section 17.12.4, “Online DDL Memory Management”](#).

[innodb_deadlock_detect](#)

Command-Line Format	--innodb-deadlock-detect[={OFF ON}]
System Variable	innodb_deadlock_detect
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	ON

This option is used to disable deadlock detection. On high concurrency systems, deadlock detection can cause a slowdown when numerous threads wait for the same lock. At times, it may be more efficient to disable deadlock detection and rely on the [innodb_lock_wait_timeout](#) setting for transaction rollback when a deadlock occurs.

For related information, see [Section 17.7.5.2, “Deadlock Detection”](#).

[innodb_dedicated_server](#)

Command-Line Format	--innodb-dedicated-server[={OFF ON}]
System Variable	innodb_dedicated_server
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	Boolean
Default Value	OFF

When [innodb_dedicated_server](#) is enabled, InnoDB automatically configures the following variables:

[innodb_buffer_pool_size](#)

[innodb_redo_log_capacity](#)

Note

[innodb_log_file_size](#) and [innodb_log_files_in_group](#) are deprecated, and superseded by [innodb_redo_log_capacity](#); see [Section 17.6.5, “Redo Log”](#).

Before MySQL 8.4, [innodb_flush_method](#) was also automatically configured.

Only consider enabling [innodb_dedicated_server](#) if the MySQL instance resides on a dedicated server where it can use all available system resources. Enabling [innodb_dedicated_server](#) is not recommended if the MySQL instance shares system resources with other applications.

For more information, see [Section 17.8.12, “Enabling Automatic Configuration for a Dedicated MySQL Server”](#).

[innodb default row format](#)

Command-Line Format	--innodb-default-row-format=value
System Variable	innodb default row format
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Enumeration
Default Value	DYNAMIC REDUNDANT
Valid Values	COMPACT DYNAMIC

The [innodb default row format](#) option defines the default row format for InnoDB tables and user-created temporary tables. The default setting is DYNAMIC. Other permitted values are COMPACT and REDUNDANT. The COMPRESSED row format, which is not supported for use in the [system tablespace](#), cannot be defined as the default.

Newly created tables use the row format defined by [innodb default row format](#) when a ROW_FORMAT option is not specified explicitly or when ROW_FORMAT=DEFAULT is used.

When a ROW_FORMAT option is not specified explicitly or when ROW_FORMAT=DEFAULT is used, any operation that rebuilds a table also silently changes the row format of the table to the format defined by [innodb default row format](#). For more information, see [Defining the Row Format of a Table](#).

Internal InnoDB temporary tables created by the server to process queries use the DYNAMIC row format, regardless of the [innodb default row format](#) setting.

[innodb directories](#)

Command-Line Format	--innodb-directories=dir_name
System Variable	innodb directories
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	Directory name
Default Value	NULL

Defines directories to scan at startup for tablespace files. This option is used when moving or restoring tablespace files to a new location while the server is offline. It is also used to specify directories of tablespace files created using an absolute path or that reside outside of the data directory.

Tablespace discovery during crash recovery relies on the [innodb directories](#) setting to identify tablespaces referenced in the redo logs. For more information, see [Tablespace Discovery During Crash Recovery](#).

The default value is NULL, but directories defined by [innodb data home dir](#), [innodb undo directory](#), and [datadir](#) are always appended to the [innodb directories](#) argument value when InnoDB builds a list of directories to scan at startup. These directories are appended regardless of whether an [innodb directories](#) setting is specified explicitly.

[innodb directories](#) may be specified as an option in a startup command or in a MySQL option file. Quotes surround the argument value because otherwise some command interpreters interpret semicolon (;) as a special character. (For example, Unix shells treat it as a command terminator.)

Startup command:

```
mysqld --innodb-directories="directory_path_1;directory_path_2"
```

MySQL option file:

```
[mysqld]  
innodb_directories="directory_path_1;directory_path_2"
```

Wildcard expressions cannot be used to specify directories.

The [innodb directories](#) scan also traverses the subdirectories of specified directories. Duplicate directories and subdirectories are discarded from the list of directories to be scanned.

For more information, see [Section 17.6.3.6, “Moving Tablespace Files While the Server is Offline”](#).

[innodb disable sort file cache](#)

Command-Line Format	--innodb-disable-sort-file-cache[={OFF ON}]
System Variable	innodb disable sort file cache
Scope	Global
Dynamic	Yes

[SET VAR](#) Hint Applies
Type
Default Value

No
 Boolean
 OFF

Disables the operating system file system cache for merge-sort temporary files. The effect is to open such files with the equivalent of O_DIRECT.

[innodb_doublewrite](#)

Command-Line Format
System Variable
Scope
Dynamic
[SET VAR](#) Hint Applies
Type
Default Value

--innodb-doublewrite=value
[innodb_doublewrite](#)
 Global
 Yes
 No
 Enumeration
 ON
 ON

Valid Values

OFF
 DETECT_AND_RECOVER
 DETECT_ONLY

The [innodb_doublewrite](#) variable controls doublewrite buffering. Doublewrite buffering is enabled by default in most cases.

You can set [innodb_doublewrite](#) to ON or OFF when starting the server to enable or disable doublewrite buffering, respectively. DETECT_AND_RECOVER is the same as ON. With this setting, except that the doublewrite buffer is fully enabled, with database page content written to the doublewrite buffer where it is accessed during recovery to fix incomplete page writes. With DETECT_ONLY, only metadata is written to the doublewrite buffer. Database page content is not written to the doublewrite buffer, and recovery does not use the doublewrite buffer to fix incomplete page writes. This lightweight setting is intended for detecting incomplete page writes only.

MySQL supports dynamic changes to the [innodb_doublewrite](#) setting that enables the doublewrite buffer, between ON, DETECT_AND_RECOVER, and DETECT_ONLY. MySQL does not support dynamic changes between a setting that enables the doublewrite buffer and OFF or vice versa.

If the doublewrite buffer is located on a Fusion-io device that supports atomic writes, the doublewrite buffer is automatically disabled and data file writes are performed using Fusion-io atomic writes instead. However, be aware that the [innodb_doublewrite](#) setting is global. When the doublewrite buffer is disabled, it is disabled for all data files including those that do not reside on Fusion-io hardware. This feature is only supported on Fusion-io hardware and is only enabled for Fusion-io NVMe on Linux. To take full advantage of this feature, an [innodb_flush_method](#) setting of O_DIRECT is recommended.

For related information, see [Section 17.6.4, “Doublewrite Buffer”](#).

[innodb_doublewrite_batch_size](#)

Command-Line Format
System Variable
Scope
Dynamic
[SET VAR](#) Hint Applies
Type
Default Value
Minimum Value
Maximum Value

--innodb-doublewrite-batch-size=#
[innodb_doublewrite_batch_size](#)
 Global
 No
 No
 Integer
 0
 0
 256

Defines the number of doublewrite pages to write in a batch.

For more information, see [Section 17.6.4, “Doublewrite Buffer”](#).

[innodb_doublewrite_dir](#)

Command-Line Format
System Variable
Scope
Dynamic
[SET VAR](#) Hint Applies
Type

--innodb-doublewrite-dir=dir_name
[innodb_doublewrite_dir](#)
 Global
 No
 No
 Directory name

Defines the directory for doublewrite files. If no directory is specified, doublewrite files are created in the [innodb_data_home_dir](#) directory, which defaults to the data directory if unspecified.

For more information, see [Section 17.6.4, “Doublewrite Buffer”](#).

[innodb_doublewrite_files](#)

Command-Line Format

--innodb-doublewrite-files=#

System Variable	innodb_doublewrite_files
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	Integer
Default Value	2
Minimum Value	1
Maximum Value	256

Defines the number of doublewrite files. By default, two doublewrite files are created for each buffer pool instance.

At a minimum, there are two doublewrite files. The maximum number of doublewrite files is two times the number of buffer pool instances. (The number of buffer pool instances is controlled by the [innodb_buffer_pool_instances](#) variable.)

For more information, see [Section 17.6.4, “Doublewrite Buffer”](#).

[innodb_doublewrite_pages](#)

Command-Line Format	--innodb-doublewrite-pages=#
System Variable	innodb_doublewrite_pages
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	Integer
Default Value	128
Minimum Value	1
Maximum Value	512

Defines the maximum number of doublewrite pages per thread for a batch write. If no value is specified, [innodb_doublewrite_pages](#) defaults to 128.

Before MySQL 8.4, the default value was the [innodb_write_io_threads](#) value, which is 4 by default.

For more information, see [Section 17.6.4, “Doublewrite Buffer”](#).

[innodb_extend_and_initialize](#)

Command-Line Format	--innodb=extend-and-initialize[={OFF ON}]
System Variable	innodb_extend_and_initialize
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	ON

Controls how space is allocated to file-per-table and general tablespaces on Linux systems.

When enabled, InnoDB writes NULLs to newly allocated pages. When disabled, space is allocated using `posix_fallocate()` calls, which reserve space without physically writing NULLs.

For more information, see [Section 17.6.3.8, “Optimizing Tablespace Space Allocation on Linux”](#).

[innodb_fast_shutdown](#)

Command-Line Format	--innodb-fast-shutdown=#
System Variable	innodb_fast_shutdown
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	1
	0
Valid Values	1
	2

The InnoDB [shutdown](#) mode. If the value is 0, InnoDB does a [slow shutdown](#), a full [purge](#) and a change buffer merge before shutting down. If the value is 1 (the default), InnoDB skips these operations at shutdown, a process known as a [fast shutdown](#). If the value is 2, InnoDB flushes its logs and shuts down cold, as if MySQL had crashed; no committed transactions are lost, but the [crash recovery](#) operation makes the next startup take longer.

The slow shutdown can take minutes, or even hours in extreme cases where substantial amounts of data are still buffered. Use the slow shutdown technique before upgrading or downgrading between MySQL major releases, so that all data files are fully prepared in case the upgrade process updates the file format.

Use [innodb_fast_shutdown=2](#) in emergency or troubleshooting situations, to get the absolute fastest shutdown if data is at risk of corruption.

[innodb_fil_make_page_dirty_debug](#)

Command-Line Format	--innodb-fil-make-page-dirty-debug=#
System Variable	innodb_fil_make_page_dirty_debug
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	0
Minimum Value	0
Maximum Value	2**32-1

By default, setting [innodb_fil_make_page_dirty_debug](#) to the ID of a tablespace immediately dirties the first page of the tablespace. If [innodb_saved_page_number_debug](#) is set to a non-default value, setting [innodb_fil_make_page_dirty_debug](#) dirties the specified page. The [innodb_fil_make_page_dirty_debug](#) option is only available if debugging support is compiled in using the [WITH_DEBUG](#) CMake option.

[innodb_file_per_table](#)

Command-Line Format	--innodb-file-per-table[={OFF ON}]
System Variable	innodb_file_per_table
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	ON

When [innodb_file_per_table](#) is enabled, tables are created in file-per-table tablespaces by default. When disabled, tables are created in the system tablespace by default. For information about file-per-table tablespaces, see [Section 17.6.3.2, “File-Per-Table Tablespaces”](#). For information about the InnoDB system tablespace, see [Section 17.6.3.1, “The System Tablespace”](#).

The [innodb_file_per_table](#) variable can be configured at runtime using a [SET GLOBAL](#) statement, specified on the command line at startup, or specified in an option file. Configuration at runtime requires privileges sufficient to set global system variables (see [Section 7.1.9.1, “System Variable Privileges”](#)) and immediately affects the operation of all connections.

When a table that resides in a file-per-table tablespace is truncated or dropped, the freed space is returned to the operating system. Truncating or dropping a table that resides in the system tablespace only frees space in the system tablespace. Freed space in the system tablespace can be used again for InnoDB data but is not returned to the operating system, as system tablespace data files never shrink.

The [innodb_file_per-table](#) setting does not affect the creation of temporary tables; temporary tables are created in session temporary tablespaces. See [Section 17.6.3.5, “Temporary Tablespaces”](#).

[innodb_fill_factor](#)

Command-Line Format	--innodb-fill-factor=#
System Variable	innodb_fill_factor
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	100
Minimum Value	10
Maximum Value	100

InnoDB performs a bulk load when creating or rebuilding indexes. This method of index creation is known as a “sorted index build”.

[innodb_fill_factor](#) defines the percentage of space on each B-tree page that is filled during a sorted index build, with the remaining space reserved for future index growth. For example, setting [innodb_fill_factor](#) to 80 reserves 20 percent of the space on each B-tree page for future index growth. Actual percentages may vary. The [innodb_fill_factor](#) setting is interpreted as a hint rather than a hard limit.

An [innodb_fill_factor](#) setting of 100 leaves 1/16 of the space in clustered index pages free for future index growth.

[innodb_fill_factor](#) applies to both B-tree leaf and non-leaf pages. It does not apply to external pages used for [TEXT](#) or [BLOB](#) entries.

For more information, see [Section 17.6.2.3, “Sorted Index Builds”](#).

[innodb_flush_log_at_timeout](#)

Command-Line Format	--innodb-flush-log-at-timeout=#
System Variable	innodb_flush_log_at_timeout
Scope	Global
Dynamic	Yes

SET VAR Hint Applies	No
Type	Integer
Default Value	1
Minimum Value	1
Maximum Value	2700
Unit	seconds

Write and flush the logs every N seconds. [innodb_flush_log_at_timeout](#) allows the timeout period between flushes to be increased in order to reduce flushing and avoid impacting performance of binary log group commit. The default setting for [innodb_flush_log_at_timeout](#) is once per second.

[innodb_flush_log_at_trx_commit](#)

Command-Line Format	--innodb-flush-log-at-trx-commit=#
System Variable	innodb_flush_log_at_trx_commit
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Enumeration
Default Value	1
	0
Valid Values	1
	2

Controls the balance between strict [ACID](#) compliance for [commit](#) operations and higher performance that is possible when commit-related I/O operations are rearranged and done in batches. You can achieve better performance by changing the default value but then you can lose transactions in a crash.

The default setting of 1 is required for full ACID compliance. Logs are written and flushed to disk at each transaction commit.

With a setting of 0, logs are written and flushed to disk once per second. Transactions for which logs have not been flushed can be lost in a crash.

With a setting of 2, logs are written after each transaction commit and flushed to disk once per second. Transactions for which logs have not been flushed can be lost in a crash.

For settings 0 and 2, once-per-second flushing is not 100% guaranteed. Flushing may occur more frequently due to DDL changes and other internal InnoDB activities that cause logs to be flushed independently of the [innodb_flush_log_at_trx_commit](#) setting, and sometimes less frequently due to scheduling issues. If logs are flushed once per second, up to one second of transactions can be lost in a crash. If logs are flushed more or less frequently than once per second, the amount of transactions that can be lost varies accordingly.

Log flushing frequency is controlled by [innodb_flush_log_at_timeout](#), which allows you to set log flushing frequency to N seconds (where N is 1 . . . 2700, with a default value of 1). However, any unexpected [mysqld](#) process exit can erase up to N seconds of transactions.

DDL changes and other internal InnoDB activities flush the log independently of the [innodb_flush_log_at_trx_commit](#) setting.

InnoDB [crash recovery](#) works regardless of the [innodb_flush_log_at_trx_commit](#) setting. Transactions are either applied entirely or erased entirely.

For durability and consistency in a replication setup that uses InnoDB with transactions:

If binary logging is enabled, set `sync_binlog=1`.

Always set [innodb_flush_log_at_trx_commit=1](#).

For information on the combination of settings on a replica that is most resilient to unexpected halts, see [Section 19.4.2, “Handling an Unexpected Halt of a Replica”](#).

Caution

Many operating systems and some disk hardware fool the flush-to-disk operation. They may tell [mysqld](#) that the flush has taken place, even though it has not. In this case, the durability of transactions is not guaranteed even with the recommended settings, and in the worst case, a power outage can corrupt InnoDB data. Using a battery-backed disk cache in the SCSI disk controller or in the disk itself speeds up file flushes, and makes the operation safer. You can also try to disable the caching of disk writes in hardware caches.

[innodb_flush_method](#)

Command-Line Format	--innodb-flush-method=value
System Variable	innodb_flush_method
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	String

Default Value (Unix)
Default Value (Windows)

O_DIRECT if supported, otherwise fsync
unbuffered
fsync

O_DSYNC

littlesync

Valid Values (Unix)

nosync

O_DIRECT

O_DIRECT_NO_FSYNC

unbuffered

Valid Values (Windows)

normal

Defines the method used to [flush](#) data to InnoDB [data files](#) and [log files](#), which can affect I/O throughput.

On Unix-like systems, the default value is O_DIRECT if supported otherwise defaults to fsync. On Windows, the default value is unbuffered.

The [innodb_flush_method](#) options for Unix-like systems include:

fsync or 0: InnoDB uses the fsync() system call to flush both the data and log files.

O_DSYNC or 1: InnoDB uses O_DSYNC to open and flush the log files, and fsync() to flush the data files. InnoDB does not use O_DSYNC directly because there have been problems with it on many varieties of Unix.

littlesync or 2: This option is used for internal performance testing and is currently unsupported. Use at your own risk.

nosync or 3: This option is used for internal performance testing and is currently unsupported. Use at your own risk.

O_DIRECT or 4: InnoDB uses O_DIRECT (or directio() on Solaris) to open the data files, and uses fsync() to flush both the data and log files. This option is available on some GNU/Linux versions, FreeBSD, and Solaris.

O_DIRECT_NO_FSYNC: InnoDB uses O_DIRECT during flushing I/O, but skips the fsync() system call after each write operation.

MySQL calls fsync() after creating a new file, after increasing file size, and after closing a file, to ensure that file system metadata changes are synchronized. The fsync() system call is still skipped after each write operation.

Data loss is possible if redo log files and data files reside on different storage devices, and an unexpected exit occurs before data file writes are flushed from a device cache that is not battery-backed. If you use or intend to use different storage devices for redo log files and data files, and your data files reside on a device with a cache that is not battery-backed, use O_DIRECT instead.

On platforms that support fdatasync() system calls, the [innodb_use_fdatasync](#) variable permits [innodb_flush_method](#) options that use fsync() to use fdatasync() instead. An fdatasync() system call does not flush changes to file metadata unless required for subsequent data retrieval, providing a potential performance benefit.

The [innodb_flush_method](#) options for Windows systems include:

unbuffered or 0: InnoDB uses non-buffered I/O.

Note

Running MySQL server on a 4K sector hard drive on Windows is not supported with unbuffered. The workaround is to use [innodb_flush_method=normal](#).

normal or 1: InnoDB uses buffered I/O.

How each setting affects performance depends on hardware configuration and workload. Benchmark your particular configuration to decide which setting to use, or whether to keep the default setting. Examine the [InnoDB_data_fsynics](#) status variable to see the overall number of fsync() calls (or fdatasync() calls if [innodb_use_fdatasync](#) is enabled) for each setting. The mix of read and write operations in your workload can affect how a setting performs. For example, on a system with a hardware RAID controller and battery-backed write cache, O_DIRECT can help to avoid double buffering between the InnoDB buffer pool and the operating system file system cache. On some systems where InnoDB data and log files are located on a SAN, the default value or O_DSYNC might be faster for a read-heavy workload with mostly SELECT statements. Always test this parameter with hardware and workload that reflect your production environment. For general I/O tuning advice, see [Section 10.5.8, "Optimizing InnoDB Disk I/O"](#).

[innodb_flush_neighbors](#)

Command-Line Format
System Variable
Scope
Dynamic

--innodb-flush-neighbors=#
[innodb_flush_neighbors](#)
Global
Yes

<u>SET VAR</u> Hint Applies	No
Type	Enumeration
Default Value	0
	0
Valid Values	1
	2

Specifies whether [flushing](#) a page from the InnoDB [buffer pool](#) also flushes other [dirty pages](#) in the same [extent](#).

A setting of 0 disables [innodb_flush_neighbors](#). Dirty pages in the same extent are not flushed.

A setting of 1 flushes contiguous dirty pages in the same extent.

A setting of 2 flushes dirty pages in the same extent.

When the table data is stored on a traditional [HDD](#) storage device, flushing such [neighbor pages](#) in one operation reduces I/O overhead (primarily for disk seek operations) compared to flushing individual pages at different times. For table data stored on [SSD](#), seek time is not a significant factor and you can set this option to 0 to spread out write operations. For related information, see [Section 17.8.3.5, “Configuring Buffer Pool Flushing”](#).

[innodb_flush_sync](#)

Command-Line Format	--innodb-flush-sync[={OFF ON}]
System Variable	<u>innodb_flush_sync</u>
Scope	Global
Dynamic	Yes
<u>SET VAR</u> Hint Applies	No
Type	Boolean
Default Value	ON

The [innodb_flush_sync](#) variable, which is enabled by default, causes the [innodb_io_capacity](#) and [innodb_io_capacity_max](#) settings to be ignored during bursts of I/O activity that occur at [checkpoints](#). To adhere to the I/O rate defined by [innodb_io_capacity](#) and [innodb_io_capacity_max](#), disable [innodb_flush_sync](#).

For information about configuring the [innodb_flush_sync](#) variable, see [Section 17.8.7, “Configuring InnoDB I/O Capacity”](#).

[innodb_flushing_avg_loops](#)

Command-Line Format	--innodb-flushing-avg-loops=#
System Variable	<u>innodb_flushing_avg_loops</u>
Scope	Global
Dynamic	Yes
<u>SET VAR</u> Hint Applies	No
Type	Integer
Default Value	30
Minimum Value	1
Maximum Value	1000

Number of iterations for which InnoDB keeps the previously calculated snapshot of the flushing state, controlling how quickly [adaptive flushing](#) responds to changing [workloads](#). Increasing the value makes the rate of [flush](#) operations change smoothly and gradually as the workload changes. Decreasing the value makes adaptive flushing adjust quickly to workload changes, which can cause spikes in flushing activity if the workload increases and decreases suddenly.

For related information, see [Section 17.8.3.5, “Configuring Buffer Pool Flushing”](#).

[innodb_force_load_corrupted](#)

Command-Line Format	--innodb-force-load-corrupted[={OFF ON}]
System Variable	<u>innodb_force_load_corrupted</u>
Scope	Global
Dynamic	No
<u>SET VAR</u> Hint Applies	No
Type	Boolean
Default Value	OFF

Permits InnoDB to load tables at startup that are marked as corrupted. Use only during troubleshooting, to recover data that is otherwise inaccessible. When troubleshooting is complete, disable this setting and restart the server.

[innodb_force_recovery](#)

Command-Line Format	--innodb-force-recovery=#
System Variable	<u>innodb_force_recovery</u>
Scope	Global

Dynamic	No
SET VAR Hint Applies	No
Type	Integer
Default Value	0
Minimum Value	0
Maximum Value	6

The [crash recovery](#) mode, typically only changed in serious troubleshooting situations. Possible values are from 0 to 6. For the meanings of these values and important information about [innodb_force_recovery](#), see [Section 17.20.3, “Forcing InnoDB Recovery”](#).

Warning

Only set this variable to a value greater than 0 in an emergency situation so that you can start InnoDB and dump your tables. As a safety measure, InnoDB prevents [INSERT](#), [UPDATE](#), or [DELETE](#) operations when [innodb_force_recovery](#) is greater than 0. An [innodb_force_recovery](#) setting of 4 or greater places InnoDB into read-only mode.

These restrictions may cause replication administration commands to fail with an error, as replication stores the replica status logs in InnoDB tables.

[innodb_fsync_threshold](#)

Command-Line Format	--innodb-fsync-threshold=#
System Variable	innodb_fsync_threshold
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	0
Minimum Value	0
Maximum Value	2**64-1

By default, when InnoDB creates a new data file, such as a new log file or tablespace file, the file is fully written to the operating system cache before it is flushed to disk, which can cause a large amount of disk write activity to occur at once. To force smaller, periodic flushes of data from the operating system cache, you can use the [innodb_fsync_threshold](#) variable to define a threshold value, in bytes. When the byte threshold is reached, the contents of the operating system cache are flushed to disk. The default value of 0 forces the default behavior, which is to flush data to disk only after a file is fully written to the cache.

Specifying a threshold to force smaller, periodic flushes may be beneficial in cases where multiple MySQL instances use the same storage devices. For example, creating a new MySQL instance and its associated data files could cause large surges of disk write activity, impeding the performance of other MySQL instances that use the same storage devices. Configuring a threshold helps avoid such surges in write activity.

[innodb_ft_aux_table](#)

System Variable	innodb_ft_aux_table
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	String

Specifies the qualified name of an InnoDB table containing a FULLTEXT index. This variable is intended for diagnostic purposes and can only be set at runtime. For example:

```
SET GLOBAL innodb_ft_aux_table = 'test/t1';
```

After you set this variable to a name in the format *db_name/table_name*, the `INFORMATION_SCHEMA` tables [INNODB_FT_INDEX_TABLE](#), [INNODB_FT_INDEX_CACHE](#), [INNODB_FT_CONFIG](#), [INNODB_FT_DELETED](#), and [INNODB_FT_BEING_DELETED](#) show information about the search index for the specified table.

For more information, see [Section 17.15.4, “InnoDB INFORMATION_SCHEMA FULLTEXT Index Tables”](#).

[innodb_ft_cache_size](#)

Command-Line Format	--innodb-ft-cache-size=#
System Variable	innodb_ft_cache_size
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	Integer
Default Value	8000000
Minimum Value	1600000
Maximum Value	80000000
Unit	bytes

The memory allocated, in bytes, for the InnoDB FULLTEXT search index cache, which holds a parsed document in memory while creating an InnoDB FULLTEXT index. Index inserts and updates are only committed to disk when the [innodb_ft_cache_size](#) size limit is reached.

[innodb_ft_cache_size](#) defines the cache size on a per table basis. To set a global limit for all tables, see [innodb_ft_total_cache_size](#).

For more information, see [InnoDB Full-Text Index Cache](#).

[innodb_ft_enable_diag_print](#)

Command-Line Format	--innodb-ft-enable-diag-print[={OFF ON}]
System Variable	innodb_ft_enable_diag_print
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	OFF

Whether to enable additional full-text search (FTS) diagnostic output. This option is primarily intended for advanced FTS debugging and is not of interest to most users. Output is printed to the error log and includes information such as:

FTS index sync progress (when the FTS cache limit is reached). For example:

```
FTS SYNC for table test, deleted count: 100 size: 10000 bytes
SYNC words: 100
```

FTS optimize progress. For example:

```
FTS start optimize test
FTS_OPTIMIZE: optimize "mysql"
FTS_OPTIMIZE: processed "mysql"
```

FTS index build progress. For example:

```
Number of doc processed: 1000
```

For FTS queries, the query parsing tree, word weight, query processing time, and memory usage are printed. For example:

```
FTS Search Processing time: 1 secs: 100 millisec: row(s) 10000
Full Search Memory: 245666 (bytes), Row: 10000
```

[innodb_ft_enable_stopword](#)

Command-Line Format	--innodb-ft-enable-stopword[={OFF ON}]
System Variable	innodb_ft_enable_stopword
Scope	Global, Session
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	ON

Specifies that a set of [stopwords](#) is associated with an InnoDB FULLTEXT index at the time the index is created. If the [innodb_ft_user_stopword_table](#) option is set, the stopwords are taken from that table. Else, if the [innodb_ft_server_stopword_table](#) option is set, the stopwords are taken from that table. Otherwise, a built-in set of default stopwords is used.

For more information, see [Section 14.9.4, "Full-Text Stopwords"](#).

[innodb_ft_max_token_size](#)

Command-Line Format	--innodb-ft-max-token-size=#
System Variable	innodb_ft_max_token_size
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	Integer
Default Value	84
Minimum Value	10
Maximum Value	84

Maximum character length of words that are stored in an InnoDB FULLTEXT index. Setting a limit on this value reduces the size of the index, thus speeding up queries, by omitting long keywords or arbitrary collections of letters that are not real words and are not likely to be search terms.

For more information, see [Section 14.9.6, "Fine-Tuning MySQL Full-Text Search"](#).

[innodb_ft_min_token_size](#)

Command-Line Format	--innodb-ft-min-token-size=#
System Variable	innodb_ft_min_token_size
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	Integer
Default Value	3
Minimum Value	0
Maximum Value	16

Minimum length of words that are stored in an InnoDB FULLTEXT index. Increasing this value reduces the size of the index, thus speeding up queries, by omitting common words that are unlikely to be significant in a search context, such as the English words “a” and “to”. For content using a CJK (Chinese, Japanese, Korean) character set, specify a value of 1.

For more information, see [Section 14.9.6, “Fine-Tuning MySQL Full-Text Search”](#).

[innodb_ft_num_word_optimize](#)

Command-Line Format	--innodb-ft-num-word-optimize=#
System Variable	innodb_ft_num_word_optimize
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	2000
Minimum Value	1000
Maximum Value	10000

Number of words to process during each [OPTIMIZE TABLE](#) operation on an InnoDB FULLTEXT index. Because a bulk insert or update operation to a table containing a full-text search index could require substantial index maintenance to incorporate all changes, you might do a series of [OPTIMIZE TABLE](#) statements, each picking up where the last left off.

For more information, see [Section 14.9.6, “Fine-Tuning MySQL Full-Text Search”](#).

[innodb_ft_result_cache_limit](#)

Command-Line Format	--innodb-ft-result-cache-limit=#
System Variable	innodb_ft_result_cache_limit
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	2000000000
Minimum Value	1000000
Maximum Value	$2^{32}-1$
Unit	bytes

The InnoDB full-text search query result cache limit (defined in bytes) per full-text search query or per thread. Intermediate and final InnoDB full-text search query results are handled in memory. Use [innodb_ft_result_cache_limit](#) to place a size limit on the full-text search query result cache to avoid excessive memory consumption in case of very large InnoDB full-text search query results (millions or hundreds of millions of rows, for example). Memory is allocated as required when a full-text search query is processed. If the result cache size limit is reached, an error is returned indicating that the query exceeds the maximum allowed memory.

The maximum value of [innodb_ft_result_cache_limit](#) for all platform types and bit sizes is $2^{32}-1$.

[innodb_ft_server_stopword_table](#)

Command-Line Format	--innodb-ft-server-stopword-table=db_name/table_name
System Variable	innodb_ft_server_stopword_table
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	String
Default Value	NULL

This option is used to specify your own InnoDB FULLTEXT index stopwords list for all InnoDB tables. To configure your own stopwords list for a specific InnoDB table, use [innodb_ft_user_stopword_table](#).

Set [innodb_ft_server_stopword_table](#) to the name of the table containing a list of stopwords, in the format *db_name/table_name*.

The stopword table must exist before you configure [innodb_ft_server_stopword_table](#). [innodb_ft_enable_stopword](#) must be enabled and [innodb_ft_server_stopword_table](#) option must be configured before you create the FULLTEXT index.

The stopword table must be an InnoDB table, containing a single VARCHAR column named value.

For more information, see [Section 14.9.4, “Full-Text Stopwords”](#).

[innodb_ft_sort_pll_degree](#)

Command-Line Format	--innodb-ft-sort-pll-degree=#
System Variable	innodb_ft_sort_pll_degree
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	Integer
Default Value	2
Minimum Value	1
Maximum Value	16

Number of threads used in parallel to index and tokenize text in an InnoDB FULLTEXT index when building a [search index](#).

For related information, see [Section 17.6.2.4, “InnoDB Full-Text Indexes”](#), and [innodb_sort_buffer_size](#).

[innodb_ft_total_cache_size](#)

Command-Line Format	--innodb-ft-total-cache-size=#
System Variable	innodb_ft_total_cache_size
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	Integer
Default Value	640000000
Minimum Value	32000000
Maximum Value	1600000000
Unit	bytes

The total memory allocated, in bytes, for the InnoDB full-text search index cache for all tables. Creating numerous tables, each with a FULLTEXT search index, could consume a significant portion of available memory. [innodb_ft_total_cache_size](#) defines a global memory limit for all full-text search indexes to help avoid excessive memory consumption. If the global limit is reached by an index operation, a forced sync is triggered.

For more information, see [InnoDB Full-Text Index Cache](#).

[innodb_ft_user_stopword_table](#)

Command-Line Format	--innodb-ft-user-stopword-table=db_name/table_name
System Variable	innodb_ft_user_stopword_table
Scope	Global, Session
Dynamic	Yes
SET VAR Hint Applies	No
Type	String
Default Value	NULL

This option is used to specify your own InnoDB FULLTEXT index stopword list on a specific table. To configure your own stopword list for all InnoDB tables, use [innodb_ft_server_stopword_table](#).

Set [innodb_ft_user_stopword_table](#) to the name of the table containing a list of stopwords, in the format *db_name/table_name*.

The stopword table must exist before you configure [innodb_ft_user_stopword_table](#). [innodb_ft_enable_stopword](#) must be enabled and [innodb_ft_user_stopword_table](#) must be configured before you create the FULLTEXT index.

The stopword table must be an InnoDB table, containing a single VARCHAR column named value.

For more information, see [Section 14.9.4, “Full-Text Stopwords”](#).

[innodb_idle_flush_pct](#)

Command-Line Format	--innodb-idle-flush-pct=#
System Variable	innodb_idle_flush_pct
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	100

Minimum Value	0
Maximum Value	100

Limits page flushing when InnoDB is idle. The [innodb_idle_flush_pct](#) value is a percentage of the [innodb_io_capacity](#) setting, which defines the number of I/O operations per second available to InnoDB. For more information, see [Limiting Buffer Flushing During Idle Periods](#).

[innodb_io_capacity](#)

Command-Line Format	--innodb-io-capacity=#
System Variable	innodb_io_capacity
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	10000
Minimum Value	100
Maximum Value (64-bit platforms, 8.4.0)	2**64-1
Maximum Value	2**32-1

The [innodb_io_capacity](#) variable defines the number of I/O operations per second (IOPS) available to InnoDB background tasks, such as [flushing](#) pages from the [buffer pool](#) and merging data from the [change buffer](#).

For information about configuring the [innodb_io_capacity](#) variable, see [Section 17.8.7, “Configuring InnoDB I/O Capacity”](#).

[innodb_io_capacity_max](#)

Command-Line Format	--innodb-io-capacity-max=#
System Variable	innodb_io_capacity_max
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	2 * innodb_io_capacity
Minimum Value	100
Maximum Value	2**32-1

If flushing activity falls behind, InnoDB can flush more aggressively, at a higher rate of I/O operations per second (IOPS) than defined by the [innodb_io_capacity](#) variable. The [innodb_io_capacity_max](#) variable defines a maximum number of IOPS performed by InnoDB background tasks in such situations. This option does not control [innodb_flush_sync](#) behavior.

The default value is twice the value of [innodb_io_capacity](#).

For information about configuring the [innodb_io_capacity_max](#) variable, see [Section 17.8.7, “Configuring InnoDB I/O Capacity”](#).

[innodb_limit_optimistic_insert_debug](#)

Command-Line Format	--innodb-limit-optimistic-insert-debug=#
System Variable	innodb_limit_optimistic_insert_debug
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	0
Minimum Value	0
Maximum Value	2**32-1

Limits the number of records per [B-tree](#) page. A default value of 0 means that no limit is imposed. This option is only available if debugging support is compiled in using the [WITH_DEBUG](#) CMake option.

[innodb_lock_wait_timeout](#)

Command-Line Format	--innodb-lock-wait-timeout=#
System Variable	innodb_lock_wait_timeout
Scope	Global, Session
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	50
Minimum Value	1
Maximum Value	1073741824
Unit	seconds

The length of time in seconds an InnoDB [transaction](#) waits for a [row lock](#) before giving up. The default value is 50 seconds. A transaction that tries to access a row that is locked by another InnoDB transaction waits at most this many seconds for write access to the row before issuing the following error:

```
ERROR 1205 (HY000): Lock wait timeout exceeded; try restarting transaction
```

When a lock wait timeout occurs, the current statement is [rolled back](#) (not the entire transaction). To have the entire transaction roll back, start the server with the `--innodb-rollback-on-timeout` option. See also [Section 17.20.5, “InnoDB Error Handling”](#).

You might decrease this value for highly interactive applications or [OLTP](#) systems, to display user feedback quickly or put the update into a queue for processing later. You might increase this value for long-running back-end operations, such as a transform step in a data warehouse that waits for other large insert or update operations to finish.

[innodb_lock_wait_timeout](#) applies to InnoDB row locks. A MySQL [table lock](#) does not happen inside InnoDB and this timeout does not apply to waits for table locks.

The lock wait timeout value does not apply to [deadlocks](#) when [innodb_deadlock_detect](#) is enabled (the default) because InnoDB detects deadlocks immediately and rolls back one of the deadlocked transactions. When [innodb_deadlock_detect](#) is disabled, InnoDB relies on [innodb_lock_wait_timeout](#) for transaction rollback when a deadlock occurs. See [Section 17.7.5.2, “Deadlock Detection”](#).

[innodb_lock_wait_timeout](#) can be set at runtime with the `SET GLOBAL` or `SET SESSION` statement. Changing the `GLOBAL` setting requires privileges sufficient to set global system variables (see [Section 7.1.9.1, “System Variable Privileges”](#)) and affects the operation of all clients that subsequently connect. Any client can change the `SESSION` setting for [innodb_lock_wait_timeout](#), which affects only that client.

[innodb_log_buffer_size](#)

Command-Line Format	<code>--innodb-log-buffer-size=#</code>
System Variable	innodb_log_buffer_size
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	67108864
Minimum Value	1048576
Maximum Value	4294967295

The size in bytes of the buffer that InnoDB uses to write to the [log files](#) on disk. The default is 64MB. A large [log buffer](#) enables large [transactions](#) to run without the need to write the log to disk before the transactions [commit](#). Thus, if you have transactions that update, insert, or delete many rows, making the log buffer larger saves disk I/O. For related information, see [Memory Configuration](#), and [Section 10.5.4, “Optimizing InnoDB Redo Logging”](#). For general I/O tuning advice, see [Section 10.5.8, “Optimizing InnoDB Disk I/O”](#).

[innodb_log_checkpoint_fuzzy_now](#)

Command-Line Format	<code>--innodb-log-checkpoint-fuzzy-now[={OFF ON}]</code>
System Variable	innodb_log_checkpoint_fuzzy_now
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	OFF

Enable this debug option to force InnoDB to write a fuzzy checkpoint. This option is only available if debugging support is compiled in using the [WITH_DEBUG CMake](#) option.

[innodb_log_checkpoint_now](#)

Command-Line Format	<code>--innodb-log-checkpoint-now[={OFF ON}]</code>
System Variable	innodb_log_checkpoint_now
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	OFF

Enable this debug option to force InnoDB to write a checkpoint. This option is only available if debugging support is compiled in using the [WITH_DEBUG CMake](#) option.

[innodb_log_checksums](#)

Command-Line Format	<code>--innodb-log-checksums[={OFF ON}]</code>
System Variable	innodb_log_checksums
Scope	Global
Dynamic	Yes

SET VAR Hint Applies**Type****Default Value**

No

Boolean

ON

Enables or disables checksums for redo log pages.

[innodb_log_checksums=ON](#) enables the CRC-32C checksum algorithm for redo log pages. When [innodb_log_checksums](#) is disabled, the contents of the redo log page checksum field are ignored.

Checksums on the redo log header page and redo log checkpoint pages are never disabled.

innodb_log_compressed_pages**Command-Line Format****System Variable****Scope****Dynamic****SET VAR Hint Applies****Type****Default Value**`--innodb-log-compressed-pages[={OFF|ON}]`[innodb_log_compressed_pages](#)

Global

Yes

No

Boolean

ON

Specifies whether images of [re-compressed pages](#) are written to the [redo log](#). Re-compression may occur when changes are made to compressed data.

[innodb_log_compressed_pages](#) is enabled by default to prevent corruption that could occur if a different version of the `zlib` compression algorithm is used during recovery. If you are certain that the `zlib` version is not subject to change, you can disable [innodb_log_compressed_pages](#) to reduce redo log generation for workloads that modify compressed data.

To measure the effect of enabling or disabling [innodb_log_compressed_pages](#), compare redo log generation for both settings under the same workload. Options for measuring redo log generation include observing the Log sequence number (LSN) in the LOG section of [SHOW ENGINE INNODB STATUS](#) output, or monitoring [Innodb os log written](#) status for the number of bytes written to the redo log files.

For related information, see [Section 17.9.1.6, “Compression for OLTP Workloads”](#).

innodb_log_file_size**Command-Line Format****Deprecated****System Variable****Scope****Dynamic****SET VAR Hint Applies****Type****Default Value****Minimum Value****Maximum Value****Unit**`--innodb-log-file-size=#`

Yes

[innodb_log_file_size](#)

Global

No

No

Integer

50331648

4194304

512GB / `innodb_log_files_in_group`

bytes

Note

[innodb_log_file_size](#) and [innodb_log_files_in_group](#) have been superseded by [innodb_redo_log_capacity](#); see [Section 17.6.5, “Redo Log”](#).

The size in bytes of each [log file](#) in a [log group](#). The combined size of log files (`innodb_log_file_size * innodb_log_files_in_group`) cannot exceed a maximum value that is slightly less than 512GB. A pair of 255 GB log files, for example, approaches the limit but does not exceed it. The default value is 48MB.

Generally, the combined size of the log files should be large enough that the server can smooth out peaks and troughs in workload activity, which often means that there is enough redo log space to handle more than an hour of write activity. The larger the value, the less checkpoint flush activity is required in the buffer pool, saving disk I/O. Larger log files also make [crash recovery](#) slower.

The minimum [innodb_log_file_size](#) is 4MB.

For related information, see [Redo Log Configuration](#). For general I/O tuning advice, see [Section 10.5.8, “Optimizing InnoDB Disk I/O”](#).

If [innodb_dedicated_server](#) is enabled, the [innodb_log_file_size](#) value is automatically configured if it is not explicitly defined. For more information, see [Section 17.8.12, “Enabling Automatic Configuration for a Dedicated MySQL Server”](#).

innodb_log_files_in_group**Command-Line Format****Deprecated****System Variable****Scope****Dynamic**`--innodb-log-files-in-group=#`

Yes

[innodb_log_files_in_group](#)

Global

No

<u>SET VAR</u> Hint Applies	No
Type	Integer
Default Value	2
Minimum Value	2
Maximum Value	100

Note

[innodb_log_file_size](#) and [innodb_log_files_in_group](#) have been superseded by [innodb_redo_log_capacity](#); see [Section 17.6.5, “Redo Log”](#).

The number of [log files](#) in the [log group](#). InnoDB writes to the files in a circular fashion. The default (and recommended) value is 2. The location of the files is specified by [innodb_log_group_home_dir](#). The combined size of log files ([innodb_log_file_size](#) * [innodb_log_files_in_group](#)) can be up to 512GB.

For related information, see [Redo Log Configuration](#).

If [innodb_dedicated_server](#) is enabled, [innodb_log_files_in_group](#) is automatically configured if it is not explicitly defined. For more information, see [Section 17.8.12, “Enabling Automatic Configuration for a Dedicated MySQL Server”](#).

[innodb_log_group_home_dir](#)

Command-Line Format	--innodb-log-group-home-dir=dir_name
System Variable	<u>innodb_log_group_home_dir</u>
Scope	Global
Dynamic	No
<u>SET VAR</u> Hint Applies	No
Type	Directory name

The directory path to the InnoDB [redo log](#) files.

For related information, see [Redo Log Configuration](#).

[innodb_log_spin_cpu_abs_lwm](#)

Command-Line Format	--innodb-log-spin-cpu-abs-lwm=#
System Variable	<u>innodb_log_spin_cpu_abs_lwm</u>
Scope	Global
Dynamic	Yes
<u>SET VAR</u> Hint Applies	No
Type	Integer
Default Value	80
Minimum Value	0
Maximum Value	4294967295

Defines the minimum amount of CPU usage below which user threads no longer spin while waiting for flushed redo. The value is expressed as a sum of CPU core usage. For example, The default value of 80 is 80% of a single CPU core. On a system with a multi-core processor, a value of 150 represents 100% usage of one CPU core plus 50% usage of a second CPU core.

For related information, see [Section 10.5.4, “Optimizing InnoDB Redo Logging”](#).

[innodb_log_spin_cpu_pct_hwm](#)

Command-Line Format	--innodb-log-spin-cpu-pct-hwm=#
System Variable	<u>innodb_log_spin_cpu_pct_hwm</u>
Scope	Global
Dynamic	Yes
<u>SET VAR</u> Hint Applies	No
Type	Integer
Default Value	50
Minimum Value	0
Maximum Value	100

Defines the maximum amount of CPU usage above which user threads no longer spin while waiting for flushed redo. The value is expressed as a percentage of the combined total processing power of all CPU cores. The default value is 50%. For example, 100% usage of two CPU cores is 50% of the combined CPU processing power on a server with four CPU cores.

The [innodb_log_spin_cpu_pct_hwm](#) variable respects processor affinity. For example, if a server has 48 cores but the [mysqld](#) process is pinned to only four CPU cores, the other 44 CPU cores are ignored.

For related information, see [Section 10.5.4, “Optimizing InnoDB Redo Logging”](#).

[innodb_log_wait_for_flush_spin_hwm](#)

Command-Line Format	--innodb-log-wait-for-flush-spin-hwm=#
System Variable	innodb_log_wait_for_flush_spin_hwm
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	400
Minimum Value	0
Maximum Value (64-bit platforms, 8.4.0)	2**64-1
Maximum Value	2**32-1
Unit	microseconds

Defines the maximum average log flush time beyond which user threads no longer spin while waiting for flushed redo. The default value is 400 microseconds.

For related information, see [Section 10.5.4, “Optimizing InnoDB Redo Logging”](#).

[innodb_log_write_ahead_size](#)

Command-Line Format	--innodb-log-write-ahead-size=#
System Variable	innodb_log_write_ahead_size
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	8192
Minimum Value	512 (log file block size)
Maximum Value	Equal to innodb_page_size
Unit	bytes

Defines the write-ahead block size for the redo log, in bytes. To avoid “read-on-write”, set [innodb_log_write_ahead_size](#) to match the operating system or file system cache block size. The default setting is 8192 bytes. Read-on-write occurs when redo log blocks are not entirely cached to the operating system or file system due to a mismatch between write-ahead block size for the redo log and operating system or file system cache block size.

Valid values for [innodb_log_write_ahead_size](#) are multiples of the InnoDB log file block size (2^n). The minimum value is the InnoDB log file block size (512). Write-ahead does not occur when the minimum value is specified. The maximum value is equal to the [innodb_page_size](#) value. If you specify a value for [innodb_log_write_ahead_size](#) that is larger than the [innodb_page_size](#) value, the [innodb_log_write_ahead_size](#) setting is truncated to the [innodb_page_size](#) value.

Setting the [innodb_log_write_ahead_size](#) value too low in relation to the operating system or file system cache block size results in “read-on-write”. Setting the value too high may have a slight impact on `fsync` performance for log file writes due to several blocks being written at once.

For related information, see [Section 10.5.4, “Optimizing InnoDB Redo Logging”](#).

[innodb_log_writer_threads](#)

Command-Line Format	--innodb-log-writer-threads[={OFF ON}]
System Variable	innodb_log_writer_threads
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	ON

Enables dedicated log writer threads for writing redo log records from the log buffer to the system buffers and flushing the system buffers to the redo log files. Dedicated log writer threads can improve performance on high-concurrency systems, but for low-concurrency systems, disabling dedicated log writer threads provides better performance.

For more information, see [Section 10.5.4, “Optimizing InnoDB Redo Logging”](#).

[innodb_lru_scan_depth](#)

Command-Line Format	--innodb-lru-scan-depth=#
System Variable	innodb_lru_scan_depth
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	1024
Minimum Value	100
Maximum Value (64-bit platforms, 8.4.0)	2**64-1
Maximum Value	2**32-1

A parameter that influences the algorithms and heuristics for the [flush](#) operation for the InnoDB [buffer pool](#). Primarily of interest to performance experts tuning I/O-intensive workloads. It specifies, per buffer pool instance, how far down the buffer pool LRU page list the page cleaner thread scans looking for [dirty pages](#) to flush. This is a background operation performed once per second.

A setting smaller than the default is generally suitable for most workloads. A value that is much higher than necessary may impact performance. Only consider increasing the value if you have spare I/O capacity under a typical workload. Conversely, if a write-intensive workload saturates your I/O capacity, decrease the value, especially in the case of a large buffer pool.

When tuning [innodb_lru_scan_depth](#), start with a low value and configure the setting upward with the goal of rarely seeing zero free pages. Also, consider adjusting [innodb_lru_scan_depth](#) when changing the number of buffer pool instances, since [innodb_lru_scan_depth](#) * [innodb_buffer_pool_instances](#) defines the amount of work performed by the page cleaner thread each second.

For related information, see [Section 17.8.3.5, “Configuring Buffer Pool Flushing”](#). For general I/O tuning advice, see [Section 10.5.8, “Optimizing InnoDB Disk I/O”](#).

[innodb_max_dirty_pages_pct](#)

Command-Line Format	--innodb-max-dirty-pages-pct=#
System Variable	innodb_max_dirty_pages_pct
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Numeric
Default Value	90
Minimum Value	0
Maximum Value	99.999

InnoDB tries to [flush](#) data from the [buffer pool](#) so that the percentage of [dirty pages](#) does not exceed this value.

The [innodb_max_dirty_pages_pct](#) setting establishes a target for flushing activity. It does not affect the rate of flushing. For information about managing the rate of flushing, see [Section 17.8.3.5, “Configuring Buffer Pool Flushing”](#).

For related information, see [Section 17.8.3.5, “Configuring Buffer Pool Flushing”](#). For general I/O tuning advice, see [Section 10.5.8, “Optimizing InnoDB Disk I/O”](#).

[innodb_max_dirty_pages_pct_lwm](#)

Command-Line Format	--innodb-max-dirty-pages-pct-lwm=#
System Variable	innodb_max_dirty_pages_pct_lwm
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Numeric
Default Value	10
Minimum Value	0
Maximum Value	99.999

Defines a low water mark representing the percentage of [dirty pages](#) at which preflushing is enabled to control the dirty page ratio. A value of 0 disables the pre-flushing behavior entirely. The configured value should always be lower than the [innodb_max_dirty_pages_pct](#) value. For more information, see [Section 17.8.3.5, “Configuring Buffer Pool Flushing”](#).

[innodb_max_purge_lag](#)

Command-Line Format	--innodb-max-purge-lag=#
System Variable	innodb_max_purge_lag
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	0
Minimum Value	0
Maximum Value	4294967295

Defines the desired maximum purge lag. If this value is exceeded, a delay is imposed on [INSERT](#), [UPDATE](#), and [DELETE](#) operations to allow time for purge to catch up. The default value is 0, which means there is no maximum purge lag and no delay.

For more information, see [Section 17.8.9, “Purge Configuration”](#).

[innodb_max_purge_lag_delay](#)

Command-Line Format	--innodb-max-purge-lag-delay=#
System Variable	innodb_max_purge_lag_delay

Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	0
Minimum Value	0
Maximum Value	10000000
Unit	microseconds

Specifies the maximum delay in microseconds for the delay imposed when the [innodb_max_purge_lag](#) threshold is exceeded. The specified [innodb_max_purge_lag_delay](#) value is an upper limit on the delay period calculated by the [innodb_max_purge_lag](#) formula.

For more information, see [Section 17.8.9, “Purge Configuration”](#).

[innodb_max_undo_log_size](#)

Command-Line Format	--innodb-max-undo-log-size=#
System Variable	innodb_max_undo_log_size
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	1073741824
Minimum Value	10485760
Maximum Value	2**64-1
Unit	bytes

Defines a threshold size for undo tablespaces. If an undo tablespace exceeds the threshold, it can be marked for truncation when [innodb_undo_log_truncate](#) is enabled. The default value is 1073741824 bytes (1024 MiB).

For more information, see [Truncating Undo Tablespaces](#).

[innodb_merge_threshold_set_all_debug](#)

Command-Line Format	--innodb-merge-threshold-set-all-debug=#
System Variable	innodb_merge_threshold_set_all_debug
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	50
Minimum Value	1
Maximum Value	50

Defines a page-full percentage value for index pages that overrides the current `MERGE_THRESHOLD` setting for all indexes that are currently in the dictionary cache. This option is only available if debugging support is compiled in using the [WITH DEBUG CMake](#) option. For related information, see [Section 17.8.11, “Configuring the Merge Threshold for Index Pages”](#).

[innodb_monitor_disable](#)

Command-Line Format	--innodb-monitor-disable={counter module pattern all}
System Variable	innodb_monitor_disable
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	String

This variable acts as a switch, disabling InnoDB [metrics counters](#). Counter data may be queried using the Information Schema [INNODB_METRICS](#) table. For usage information, see [Section 17.15.6, “InnoDB INFORMATION_SCHEMA Metrics Table”](#).

[innodb_monitor_disable='latch'](#) disables statistics collection for [SHOW ENGINE INNODB MUTEX](#). For more information, see [Section 15.7.7.16, “SHOW ENGINE Statement”](#).

[innodb_monitor_enable](#)

Command-Line Format	--innodb-monitor-enable={counter module pattern all}
System Variable	innodb_monitor_enable
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	String

This variable acts as a switch, enabling InnoDB [metrics counters](#). Counter data may be queried using the Information Schema [INNODB_METRICS](#) table. For usage information, see [Section 17.15.6, “InnoDB INFORMATION_SCHEMA Metrics Table”](#).

[innodb_monitor_enable='latch'](#) enables statistics collection for [SHOW ENGINE INNODB MUTEX](#). For more information, see [Section 15.7.7.16, “SHOW ENGINE Statement”](#).

[innodb_monitor_reset](#)

Command-Line Format	--innodb-monitor-reset={counter module pattern all}
System Variable	innodb_monitor_reset
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Enumeration
Default Value	NULL
	counter
	module
Valid Values	pattern
	all

This variable acts as a switch, resetting the count value for InnoDB [metrics counters](#) to zero. Counter data may be queried using the Information Schema [INNODB_METRICS](#) table. For usage information, see [Section 17.15.6, “InnoDB INFORMATION_SCHEMA Metrics Table”](#).

[innodb_monitor_reset='latch'](#) resets statistics reported by [SHOW ENGINE INNODB MUTEX](#). For more information, see [Section 15.7.7.16, “SHOW ENGINE Statement”](#).

[innodb_monitor_reset_all](#)

Command-Line Format	--innodb-monitor-reset-all={counter module pattern all}
System Variable	innodb_monitor_reset_all
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Enumeration
Default Value	NULL
	counter
	module
Valid Values	pattern
	all

This variable acts as a switch, resetting all values (minimum, maximum, and so on) for InnoDB [metrics counters](#). Counter data may be queried using the Information Schema [INNODB_METRICS](#) table. For usage information, see [Section 17.15.6, “InnoDB INFORMATION_SCHEMA Metrics Table”](#).

[innodb_numa_interleave](#)

Command-Line Format	--innodb-numa-interleave[={OFF ON}]
System Variable	innodb_numa_interleave
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	Boolean
Default Value	ON

Enables the NUMA interleave memory policy for allocation of the InnoDB buffer pool. When [innodb_numa_interleave](#) is enabled, the NUMA memory policy is set to MPOL_INTERLEAVE for the [mysqld](#) process. After the InnoDB buffer pool is allocated, the NUMA memory policy is set back to MPOL_DEFAULT. For the [innodb_numa_interleave](#) option to be available, MySQL must be compiled on a NUMA-enabled Linux system. The default value is ON if the system supports it, otherwise it defaults to OFF.

CMake sets the default [WITH_NUMA](#) value based on whether the current platform has NUMA support. For more information, see [Section 2.8.7, “MySQL Source-Configuration Options”](#).

[innodb_old_blocks_pct](#)

Command-Line Format	--innodb-old-blocks-pct=#
System Variable	innodb_old_blocks_pct
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No

Type	Integer
Default Value	37
Minimum Value	5
Maximum Value	95

Specifies the approximate percentage of the InnoDB [buffer pool](#) used for the old block [sublist](#). The range of values is 5 to 95. The default value is 37 (that is, 3/8 of the pool). Often used in combination with [innodb_old_blocks_time](#).

For more information, see [Section 17.8.3.3, “Making the Buffer Pool Scan Resistant”](#). For information about buffer pool management, the [LRU](#) algorithm, and [eviction](#) policies, see [Section 17.5.1, “Buffer Pool”](#).

[innodb_old_blocks_time](#)

Command-Line Format	--innodb-old-blocks-time=#
System Variable	innodb_old_blocks_time
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	1000
Minimum Value	0
Maximum Value	2**32-1
Unit	milliseconds

Non-zero values protect against the [buffer pool](#) being filled by data that is referenced only for a brief period, such as during a [full table scan](#). Increasing this value offers more protection against full table scans interfering with data cached in the buffer pool.

Specifies how long in milliseconds a block inserted into the old [sublist](#) must stay there after its first access before it can be moved to the new sublist. If the value is 0, a block inserted into the old sublist moves immediately to the new sublist the first time it is accessed, no matter how soon after insertion the access occurs. If the value is greater than 0, blocks remain in the old sublist until an access occurs at least that many milliseconds after the first access. For example, a value of 1000 causes blocks to stay in the old sublist for 1 second after the first access before they become eligible to move to the new sublist.

The default value is 1000.

This variable is often used in combination with [innodb_old_blocks_pct](#). For more information, see [Section 17.8.3.3, “Making the Buffer Pool Scan Resistant”](#). For information about buffer pool management, the [LRU](#) algorithm, and [eviction](#) policies, see [Section 17.5.1, “Buffer Pool”](#).

[innodb_online_alter_log_max_size](#)

Command-Line Format	--innodb-online-alter-log-max-size=#
System Variable	innodb_online_alter_log_max_size
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	134217728
Minimum Value	65536
Maximum Value	2**64-1
Unit	bytes

Specifies an upper limit in bytes on the size of the temporary log files used during [online DDL](#) operations for InnoDB tables. There is one such log file for each index being created or table being altered. This log file stores data inserted, updated, or deleted in the table during the DDL operation. The temporary log file is extended when needed by the value of [innodb_sort_buffer_size](#), up to the maximum specified by [innodb_online_alter_log_max_size](#). If a temporary log file exceeds the upper size limit, the [ALTER TABLE](#) operation fails and all uncommitted concurrent DML operations are rolled back. Thus, a large value for this option allows more DML to happen during an online DDL operation, but also extends the period of time at the end of the DDL operation when the table is locked to apply the data from the log.

[innodb_open_files](#)

Command-Line Format	--innodb-open-files=#
System Variable	innodb_open_files
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	-1 (signifies autosizing; do not assign this literal value)
Minimum Value	10
Maximum Value	2147483647

Specifies the maximum number of files that InnoDB can have open at one time. The minimum value is 10. If [innodb_file_per_table](#) is disabled, the default value is 300; otherwise, the default value is 300 or the [table_open_cache](#) setting, whichever is higher.

The [innodb_open_files](#) limit can be set at runtime using a `SELECT innodb_set_open_files_limit(N)` statement, where *N* is the desired [innodb_open_files](#) limit; for example:

```
mysql> SELECT innodb_set_open_files_limit(1000);
```

The statement executes a stored procedure that sets the new limit. If the procedure is successful, it returns the value of the newly set limit; otherwise, a failure message is returned.

It is not permitted to set [innodb_open_files](#) using a `SET` statement. To set [innodb_open_files](#) at runtime, use the `SELECT innodb_set_open_files_limit(N)` statement described above.

Setting [innodb_open_files=default](#) is not supported. Only integer values are permitted.

To prevent non-LRU managed files from consuming the entire [innodb_open_files](#) limit, non-LRU managed files are limited to 90 percent of this limit, which reserves 10 percent of it for LRU managed files.

[innodb_optimize_fulltext_only](#)

Command-Line Format	<code>--innodb-optimize-fulltext-only[={OFF ON}]</code>
System Variable	innodb_optimize_fulltext_only
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	OFF

Changes the way [OPTIMIZE TABLE](#) operates on InnoDB tables. Intended to be enabled temporarily, during maintenance operations for InnoDB tables with FULLTEXT indexes.

By default, [OPTIMIZE TABLE](#) reorganizes data in the [clustered index](#) of the table. When this option is enabled, [OPTIMIZE TABLE](#) skips the reorganization of table data, and instead processes newly added, deleted, and updated token data for InnoDB FULLTEXT indexes. For more information, see [Optimizing InnoDB Full-Text Indexes](#).

[innodb_page_cleaners](#)

Command-Line Format	<code>--innodb-page-cleaners=#</code>
System Variable	innodb_page_cleaners
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	Integer
Default Value	innodb_buffer_pool_instances
Minimum Value	1
Maximum Value	64

The number of page cleaner threads that flush dirty pages from buffer pool instances. Page cleaner threads perform flush list and LRU flushing. When there are multiple page cleaner threads, buffer pool flushing tasks for each buffer pool instance are dispatched to idle page cleaner threads. The [innodb_page_cleaners](#) default value is set to the same value as [innodb_buffer_pool_instances](#). If the specified number of page cleaner threads exceeds the number of buffer pool instances, then [innodb_page_cleaners](#) is automatically set to the same value as [innodb_buffer_pool_instances](#).

If your workload is write-IO bound when flushing dirty pages from buffer pool instances to data files, and if your system hardware has available capacity, increasing the number of page cleaner threads may help improve write-IO throughput.

Multithreaded page cleaner support extends to shutdown and recovery phases.

The `setpriority()` system call is used on Linux platforms where it is supported, and where the [mysqld](#) execution user is authorized to give `page_cleaner` threads priority over other MySQL and InnoDB threads to help page flushing keep pace with the current workload. `setpriority()` support is indicated by this InnoDB startup message:

```
[Note] InnoDB: If the mysqld execution user is authorized, page cleaner thread priority can be changed. See the man page of setpriority().
```

For systems where server startup and shutdown is not managed by systemd, [mysqld](#) execution user authorization can be configured in `/etc/security/limits.conf`. For example, if [mysqld](#) is run under the `mysql` user, you can authorize the `mysql` user by adding these lines to `/etc/security/limits.conf`:

```
mysql      hard    nice    -20
mysql      soft    nice    -20
```

For systemd managed systems, the same can be achieved by specifying `LimitNICE=-20` in a localized systemd configuration file. For example, create a file named `override.conf` in `/etc/systemd/system/mysqld.service.d/override.conf` and add this entry:

```
[Service]
LimitNICE=-20
```

After creating or changing `override.conf`, reload the systemd configuration, then tell systemd to restart the MySQL service:

```
systemctl daemon-reload
systemctl restart mysqld # RPM platforms
systemctl restart mysql  # Debian platforms
```

For more information about using a localized systemd configuration file, see [Configuring systemd for MySQL](#).

After authorizing the `mysqld` execution user, use the `cat` command to verify the configured `Nice` limits for the `mysqld` process:

```
$> cat /proc/mysqld_pid/limits | grep nice
Max nice priority      18446744073709551596 18446744073709551596
```

[innodb_page_size](#)

Command-Line Format	--innodb-page-size=#
System Variable	innodb_page_size
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	Enumeration
Default Value	16384
	4096
	8192
Valid Values	16384
	32768
	65536

Specifies the [page size](#) for InnoDB [tablespaces](#). Values can be specified in bytes or kilobytes. For example, a 16 kilobyte page size value can be specified as 16384, 16KB, or 16k.

[innodb_page_size](#) can only be configured prior to initializing the MySQL instance and cannot be changed afterward. If no value is specified, the instance is initialized using the default page size. See [Section 17.8.1, “InnoDB Startup Configuration”](#).

For both 32KB and 64KB page sizes, the maximum row length is approximately 16000 bytes. `ROW_FORMAT=COMPRESSED` is not supported when [innodb_page_size](#) is set to 32KB or 64KB. For [innodb_page_size=32KB](#), extent size is 2MB. For [innodb_page_size=64KB](#), extent size is 4MB. [innodb_log_buffer_size](#) should be set to at least 16MB (the default is 64MB) when using 32KB or 64KB page sizes.

The default 16KB page size or larger is appropriate for a wide range of [workloads](#), particularly for queries involving table scans and DML operations involving bulk updates. Smaller page sizes might be more efficient for [OLTP](#) workloads involving many small writes, where contention can be an issue when single pages contain many rows. Smaller pages might also be efficient with [SSD](#) storage devices, which typically use small block sizes. Keeping the InnoDB page size close to the storage device block size minimizes the amount of unchanged data that is rewritten to disk.

The minimum file size for the first system tablespace data file (`ibdata1`) differs depending on the [innodb_page_size](#) value. See the [innodb_data_file_path](#) option description for more information.

A MySQL instance using a particular InnoDB page size cannot use data files or log files from an instance that uses a different page size.

For general I/O tuning advice, see [Section 10.5.8, “Optimizing InnoDB Disk I/O”](#).

[innodb_parallel_read_threads](#)

Command-Line Format	--innodb-parallel-read-threads=#
System Variable	innodb_parallel_read_threads
Scope	Session
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	(available logical processors / 8), min of 4
Minimum Value	1
Maximum Value	256

Defines the number of threads that can be used for parallel clustered index reads. Parallel scanning of partitions is also supported. Parallel read threads can improve [CHECK TABLE](#) performance. InnoDB reads the clustered index twice during a [CHECK TABLE](#) operation. The second read can be performed in parallel. This feature does not apply to secondary index scans. The [innodb_parallel_read_threads](#) session variable must be set to a value greater than 1 for parallel clustered index reads to occur. The actual number of threads used to perform a parallel clustered index read is determined by the

[innodb_parallel_read_threads](#) setting or the number of index subtrees to scan, whichever is smaller. The pages read into the buffer pool during the scan are kept at the tail of the buffer pool LRU list so that they can be discarded quickly when free buffer pool pages are required.

The maximum number of parallel read threads (256) is the total number of threads for all client connections. If the thread limit is reached, connections fall back to using a single thread. The default value is calculated by the number of available logical processors on the system divided by 8, with a minimum default value of 4.

Before MySQL 8.4, the default value was always 4.

[innodb_print_all_deadlocks](#)

Command-Line Format	--innodb-print-all-deadlocks[={OFF ON}]
System Variable	innodb_print_all_deadlocks
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	OFF

When this option is enabled, information about all [deadlocks](#) in InnoDB user transactions is recorded in the `mysqld` [error log](#). Otherwise, you see information about only the last deadlock, using the `SHOW ENGINE INNODB STATUS` statement. An occasional InnoDB deadlock is not necessarily an issue, because InnoDB detects the condition immediately and rolls back one of the transactions automatically. You might use this option to troubleshoot why deadlocks are occurring if an application does not have appropriate error-handling logic to detect the rollback and retry its operation. A large number of deadlocks might indicate the need to restructure transactions that issue [DML](#) or `SELECT ... FOR UPDATE` statements for multiple tables, so that each transaction accesses the tables in the same order, thus avoiding the deadlock condition.

For related information, see [Section 17.7.5, “Deadlocks in InnoDB”](#).

[innodb_print_ddl_logs](#)

Command-Line Format	--innodb-print-ddl-logs[={OFF ON}]
System Variable	innodb_print_ddl_logs
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	OFF

Enabling this option causes MySQL to write DDL logs to `stderr`. For more information, see [Viewing DDL Logs](#).

[innodb_purge_batch_size](#)

Command-Line Format	--innodb-purge-batch-size=#
System Variable	innodb_purge_batch_size
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	300
Minimum Value	1
Maximum Value	5000

Defines the number of undo log pages that purge parses and processes in one batch from the [history list](#). In a multithreaded purge configuration, the coordinator purge thread divides [innodb_purge_batch_size](#) by [innodb_purge_threads](#) and assigns that number of pages to each purge thread. The [innodb_purge_batch_size](#) variable also defines the number of undo log pages that purge frees after every 128 iterations through the undo logs.

The [innodb_purge_batch_size](#) option is intended for advanced performance tuning in combination with the [innodb_purge_threads](#) setting. Most users need not change [innodb_purge_batch_size](#) from its default value.

For related information, see [Section 17.8.9, “Purge Configuration”](#).

[innodb_purge_threads](#)

Command-Line Format	--innodb-purge-threads=#
System Variable	innodb_purge_threads
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	Integer
Default Value	1 if # of available logical processors is <= 16; otherwise 4
Minimum Value	1

The number of background threads devoted to the InnoDB [purge](#) operation. Increasing the value creates additional purge threads, which can improve efficiency on systems where [DML](#) operations are performed on multiple tables.

For related information, see [Section 17.8.9, “Purge Configuration”](#).

[innodb_purge_rseg_truncate_frequency](#)

Command-Line Format	<code>--innodb-purge-rseg-truncate-frequency=#</code>
System Variable	innodb_purge_rseg_truncate_frequency
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	128
Minimum Value	1
Maximum Value	128

Defines the frequency with which the purge system frees rollback segments in terms of the number of times that purge is invoked. An undo tablespace cannot be truncated until its rollback segments are freed. Normally, the purge system frees rollback segments once every 128 times that purge is invoked. The default value is 128. Reducing this value increases the frequency with which the purge thread frees rollback segments.

[innodb_purge_rseg_truncate_frequency](#) is intended for use with [innodb_undo_log_truncate](#). For more information, see [Truncating Undo Tablespaces](#).

[innodb_random_read_ahead](#)

Command-Line Format	<code>--innodb-random-read-ahead[={OFF ON}]</code>
System Variable	innodb_random_read_ahead
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	OFF

Enables the random [read-ahead](#) technique for optimizing InnoDB I/O.

For details about performance considerations for different types of read-ahead requests, see [Section 17.8.3.4, “Configuring InnoDB Buffer Pool Prefetching \(Read-Ahead\)”](#). For general I/O tuning advice, see [Section 10.5.8, “Optimizing InnoDB Disk I/O”](#).

[innodb_read_ahead_threshold](#)

Command-Line Format	<code>--innodb-read-ahead-threshold=#</code>
System Variable	innodb_read_ahead_threshold
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	56
Minimum Value	0
Maximum Value	64

Controls the sensitivity of linear [read-ahead](#) that InnoDB uses to prefetch pages into the [buffer pool](#). If InnoDB reads at least [innodb_read_ahead_threshold](#) pages sequentially from an [extent](#) (64 pages), it initiates an asynchronous read for the entire following extent. The permissible range of values is 0 to 64. A value of 0 disables read-ahead. For the default of 56, InnoDB must read at least 56 pages sequentially from an extent to initiate an asynchronous read for the following extent.

Knowing how many pages are read through the read-ahead mechanism, and how many of these pages are evicted from the buffer pool without ever being accessed, can be useful when fine-tuning the [innodb_read_ahead_threshold](#) setting. `SHOW ENGINE INNODB STATUS` output displays counter information from the [innodb_buffer_pool_read_ahead](#) and [innodb_buffer_pool_read_ahead_evicted](#) global status variables, which report the number of pages brought into the [buffer pool](#) by read-ahead requests, and the number of such pages [evicted](#) from the buffer pool without ever being accessed, respectively. The status variables report global values since the last server restart.

`SHOW ENGINE INNODB STATUS` also shows the rate at which the read-ahead pages are read and the rate at which such pages are evicted without being accessed. The per-second averages are based on the statistics collected since the last invocation of `SHOW ENGINE INNODB STATUS` and are displayed in the BUFFER POOL AND MEMORY section of the `SHOW ENGINE INNODB STATUS` output.

For more information, see [Section 17.8.3.4, “Configuring InnoDB Buffer Pool Prefetching \(Read-Ahead\)”](#). For general I/O tuning advice, see [Section 10.5.8, “Optimizing InnoDB Disk I/O”](#).

[innodb_read_io_threads](#)

Command-Line Format	--innodb-read-io-threads=#
System Variable	innodb_read_io_threads
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	Integer
Default Value	(available logical processors / 2), min of 4
Minimum Value	1
Maximum Value	64

The number of I/O threads for read operations in InnoDB. Its counterpart for write threads is [innodb_write_io_threads](#). For more information, see [Section 17.8.5, “Configuring the Number of Background InnoDB I/O Threads”](#). For general I/O tuning advice, see [Section 10.5.8, “Optimizing InnoDB Disk I/O”](#). The default value is the number of available logical processors on the system divided by 2, with a minimum default value of 4.

Before MySQL 8.4, the default value was always 4.

Note

On Linux systems, running multiple MySQL servers (typically more than 12) with default settings for [innodb_read_io_threads](#), [innodb_write_io_threads](#), and the Linux `aio-max-nr` setting can exceed system limits. Ideally, increase the `aio-max-nr` setting; as a workaround, you might reduce the settings for one or both of the MySQL variables.

[innodb_read_only](#)

Command-Line Format	--innodb-read-only[={OFF ON}]
System Variable	innodb_read_only
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	Boolean
Default Value	OFF

Starts InnoDB in read-only mode. For distributing database applications or data sets on read-only media. Can also be used in data warehouses to share the same data directory between multiple instances. For more information, see [Section 17.8.2, “Configuring InnoDB for Read-Only Operation”](#).

Enabling [innodb_read_only](#) prevents creating and dropping tables for all storage engines, and not only InnoDB. Table creation and drop operations for any storage engine modify data dictionary tables in the `mysql` system database, but those tables use the InnoDB storage engine and cannot be modified when [innodb_read_only](#) is enabled. The same principle applies to other table operations that require modifying data dictionary tables. Examples:

If the [innodb_read_only](#) system variable is enabled, [ANALYZE TABLE](#) may fail because it cannot update statistics tables in the data dictionary, which use InnoDB. For [ANALYZE TABLE](#) operations that update the key distribution, failure may occur even if the operation updates the table itself (for example, if it is a MyISAM table). To obtain the updated distribution statistics, set [information_schema_stats_expiry=0](#).

[ALTER TABLE tbl_name ENGINE=engine_name](#) fails because it updates the storage engine designation, which is stored in the data dictionary.

In addition, other tables in the `mysql` system database use the InnoDB storage engine. Making those tables read-only results in restrictions on operations that modify them. Examples:

Account-management statements such as [CREATE USER](#) and [GRANT](#) fail because the grant tables use InnoDB.

The [INSTALL PLUGIN](#) and [UNINSTALL PLUGIN](#) plugin-management statements fail because the `mysql.plugin` system table uses InnoDB.

The [CREATE FUNCTION](#) and [DROP FUNCTION](#) loadable function-management statements fail because the `mysql.func` system table uses InnoDB.

[innodb_redo_log_archive_dirs](#)

Command-Line Format	--innodb-redo-log-archive-dirs
System Variable	innodb_redo_log_archive_dirs
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	String
Default Value	NULL

Defines labeled directories where redo log archive files can be created. You can define multiple labeled directories in a semicolon-separated list. For example:

```
innodb_redo_log_archive_dirs='label1:/backups1;label2:/backups2'
```

A label can be any string of characters, with the exception of colons (:), which are not permitted. An empty label is also permitted, but the colon (:) is still required in this case.

A path must be specified, and the directory must exist. The path can contain colons (':'), but semicolons (;) are not permitted.

[innodb_redo_log_capacity](#)

Command-Line Format	--innodb-redo-log-capacity=#
System Variable	innodb_redo_log_capacity
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	104857600
Minimum Value	8388608
Maximum Value	549755813888
Unit	bytes

Defines the amount of disk space occupied by redo log files.

innodb_redo_log_capacity supercedes the [innodb_log_files_in_group](#) and [innodb_log_file_size](#) variables, which are both ignored if innodb_redo_log_capacity is defined.

If innodb_redo_log_capacity is not defined, and if neither innodb_log_file_size or innodb_log_files_in_group are defined, then the default innodb_redo_log_capacity value is used.

If innodb_redo_log_capacity is not defined, and if innodb_log_file_size and/or innodb_log_files_in_group is defined, then the InnoDB redo log capacity is calculated as $(innodb_log_files_in_group * innodb_log_file_size)$. This calculation does not modify the unused innodb_redo_log_capacity setting's value.

The [InnoDB_redo_log_capacity_resized](#) server status variable indicates the total redo log capacity for all redo log files.

If [innodb_dedicated_server](#) is enabled, the [innodb_redo_log_capacity](#) value is automatically configured if it is not explicitly defined. For more information, see [Section 17.8.12, “Enabling Automatic Configuration for a Dedicated MySQL Server”](#).

For more information, see [Section 17.6.5, “Redo Log”](#).

[innodb_redo_log_encrypt](#)

Command-Line Format	--innodb-redo-log-encrypt[={OFF ON}]
System Variable	innodb_redo_log_encrypt
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	OFF

Controls encryption of redo log data for tables encrypted using the InnoDB [data-at-rest encryption feature](#). Encryption of redo log data is disabled by default. For more information, see [Redo Log Encryption](#).

[innodb_replication_delay](#)

Command-Line Format	--innodb-replication-delay=#
System Variable	innodb_replication_delay
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	0
Minimum Value	0
Maximum Value	4294967295
Unit	milliseconds

The replication thread delay in milliseconds on a replica server if [innodb_thread_concurrency](#) is reached.

[innodb_rollback_on_timeout](#)

Command-Line Format	--innodb-rollback-on-timeout[={OFF ON}]
System Variable	innodb_rollback_on_timeout
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	Boolean
Default Value	OFF

InnoDB [rolls back](#) only the last statement on a transaction timeout by default. If [--innodb-rollback-on-timeout](#) is specified, a transaction timeout causes InnoDB to abort and roll back the entire transaction.

For more information, see [Section 17.20.5, “InnoDB Error Handling”](#).

[innodb_rollback_segments](#)

Command-Line Format	<code>--innodb-rollback-segments=#</code>
System Variable	innodb_rollback_segments
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	128
Minimum Value	1
Maximum Value	128

[innodb_rollback_segments](#) defines the number of [rollback segments](#) allocated to each undo tablespace and the global temporary tablespace for transactions that generate undo records. The number of transactions that each rollback segment supports depends on the InnoDB page size and the number of undo logs assigned to each transaction. For more information, see [Section 17.6.6, “Undo Logs”](#).

For related information, see [Section 17.3, “InnoDB Multi-Versioning”](#). For information about undo tablespaces, see [Section 17.6.3.4, “Undo Tablespaces”](#).

[innodb_saved_page_number_debug](#)

Command-Line Format	<code>--innodb-saved-page-number-debug=#</code>
System Variable	innodb_saved_page_number_debug
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	0
Minimum Value	0
Maximum Value	$2^{32}-1$

Saves a page number. Setting the [innodb_fil_make_page_dirty_debug](#) option dirties the page defined by [innodb_saved_page_number_debug](#). The [innodb_saved_page_number_debug](#) option is only available if debugging support is compiled in using the [WITH_DEBUG CMake](#) option.

[innodb_segment_reserve_factor](#)

Command-Line Format	<code>--innodb-segment-reserve-factor=#</code>
System Variable	innodb_segment_reserve_factor
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Numeric
Default Value	12.5
Minimum Value	0.03
Maximum Value	40

Defines the percentage of tablespace file segment pages reserved as empty pages. The setting is applicable to file-per-table and general tablespaces. The [innodb_segment_reserve_factor](#) default setting is 12.5 percent, which is the same percentage of pages reserved in previous MySQL releases.

For more information, see [Configuring the Percentage of Reserved File Segment Pages](#).

[innodb_sort_buffer_size](#)

Command-Line Format	<code>--innodb-sort-buffer-size=#</code>
System Variable	innodb_sort_buffer_size
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	Integer
Default Value	1048576
Minimum Value	65536
Maximum Value	67108864
Unit	bytes

This variable defines the amount by which the temporary log file is extended when recording concurrent DML during an [online DDL](#) operation, and the size of the temporary log file read buffer and write buffer.

For more information, see [Section 17.12.3, “Online DDL Space Requirements”](#).

[innodb_spin_wait_delay](#)

Command-Line Format	--innodb-spin-wait-delay=#
System Variable	innodb_spin_wait_delay
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	6
Minimum Value	0
Maximum Value	1000

The maximum delay between polls for a [spin](#) lock. The low-level implementation of this mechanism varies depending on the combination of hardware and operating system, so the delay does not correspond to a fixed time interval.

Can be used in combination with the [innodb_spin_wait_pause_multiplier](#) variable for greater control over the duration of spin-lock polling delays.

For more information, see [Section 17.8.8, “Configuring Spin Lock Polling”](#).

[innodb_spin_wait_pause_multiplier](#)

Command-Line Format	--innodb-spin-wait-pause-multiplier=#
System Variable	innodb_spin_wait_pause_multiplier
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	50
Minimum Value	0
Maximum Value	100

Defines a multiplier value used to determine the number of PAUSE instructions in spin-wait loops that occur when a thread waits to acquire a mutex or rw-lock.

For more information, see [Section 17.8.8, “Configuring Spin Lock Polling”](#).

[innodb_stats_auto_recalc](#)

Command-Line Format	--innodb-stats-auto-recalc[={OFF ON}]
System Variable	innodb_stats_auto_recalc
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	ON

Causes InnoDB to automatically recalculate [persistent statistics](#) after the data in a table is changed substantially. The threshold value is 10% of the rows in the table. This setting applies to tables created when the [innodb_stats_persistent](#) option is enabled. Automatic statistics recalculation may also be configured by specifying STATS_AUTO_RECALC=1 in a [CREATE TABLE](#) or [ALTER TABLE](#) statement. The amount of data sampled to produce the statistics is controlled by the [innodb_stats_persistent_sample_pages](#) variable.

For more information, see [Section 17.8.10.1, “Configuring Persistent Optimizer Statistics Parameters”](#).

[innodb_stats_include_delete_marked](#)

Command-Line Format	--innodb-stats-include-delete-marked[={OFF ON}]
System Variable	innodb_stats_include_delete_marked
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	OFF

By default, InnoDB reads uncommitted data when calculating statistics. In the case of an uncommitted transaction that deletes rows from a table, InnoDB excludes records that are delete-marked when calculating row estimates and index statistics, which can lead to non-optimal execution plans for other transactions that are operating on the table concurrently using a transaction isolation level other than [READ UNCOMMITTED](#). To avoid this scenario, [innodb_stats_include_delete_marked](#) can be enabled to ensure that InnoDB includes delete-marked records when calculating persistent optimizer statistics.

When [innodb_stats_include_delete_marked](#) is enabled, [ANALYZE TABLE](#) considers delete-marked records when recalculating statistics.

[innodb_stats_include_delete_marked](#) is a global setting that affects all InnoDB tables. It is only applicable to persistent optimizer statistics.

For related information, see [Section 17.8.10.1, “Configuring Persistent Optimizer Statistics Parameters”](#).

[innodb_stats_method](#)

Command-Line Format	<code>--innodb-stats-method=value</code>
System Variable	innodb_stats_method
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Enumeration
Default Value	<code>nulls_equal</code> <code>nulls_equal</code>
Valid Values	<code>nulls_unequal</code> <code>nulls_ignored</code>

How the server treats NULL values when collecting [statistics](#) about the distribution of index values for InnoDB tables. Permitted values are `nulls_equal`, `nulls_unequal`, and `nulls_ignored`. For `nulls_equal`, all NULL index values are considered equal and form a single value group with a size equal to the number of NULL values. For `nulls_unequal`, NULL values are considered unequal, and each NULL forms a distinct value group of size 1. For `nulls_ignored`, NULL values are ignored.

The method used to generate table statistics influences how the optimizer chooses indexes for query execution, as described in [Section 10.3.8, “InnoDB and MyISAM Index Statistics Collection”](#).

[innodb_stats_on_metadata](#)

Command-Line Format	<code>--innodb-stats-on-metadata[={OFF ON}]</code>
System Variable	innodb_stats_on_metadata
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	OFF

This option only applies when optimizer [statistics](#) are configured to be non-persistent. Optimizer statistics are not persisted to disk when [innodb_stats_persistent](#) is disabled or when individual tables are created or altered with `STATS_PERSISTENT=0`. For more information, see [Section 17.8.10.2, “Configuring Non-Persistent Optimizer Statistics Parameters”](#).

When [innodb_stats_on_metadata](#) is enabled, InnoDB updates non-persistent [statistics](#) when metadata statements such as [SHOW TABLE STATUS](#) or when accessing the Information Schema [TABLES](#) or [STATISTICS](#) tables. (These updates are similar to what happens for [ANALYZE TABLE](#).) When disabled, InnoDB does not update statistics during these operations. Leaving the setting disabled can improve access speed for schemas that have a large number of tables or indexes. It can also improve the stability of [execution plans](#) for queries that involve InnoDB tables.

To change the setting, issue the statement `SET GLOBAL innodb_stats_on_metadata=mode`, where *mode* is either ON or OFF (or 1 or 0). Changing the setting requires privileges sufficient to set global system variables (see [Section 7.1.9.1, “System Variable Privileges”](#)) and immediately affects the operation of all connections.

[innodb_stats_persistent](#)

Command-Line Format	<code>--innodb-stats-persistent[={OFF ON}]</code>
System Variable	innodb_stats_persistent
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	ON

Specifies whether InnoDB index statistics are persisted to disk. Otherwise, statistics may be recalculated frequently which can lead to variations in [query execution plans](#). This setting is stored with each table when the table is created. You can set [innodb_stats_persistent](#) at the global level before creating a table, or use the `STATS_PERSISTENT` clause of the [CREATE TABLE](#) and [ALTER TABLE](#) statements to override the system-wide setting and configure persistent statistics for individual tables.

For more information, see [Section 17.8.10.1, “Configuring Persistent Optimizer Statistics Parameters”](#).

[innodb_stats_persistent_sample_pages](#)

Command-Line Format	<code>--innodb-stats-persistent-sample-pages=#</code>
System Variable	innodb_stats_persistent_sample_pages
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer

Default Value	20
Minimum Value	1
Maximum Value	18446744073709551615

The number of index [pages](#) to sample when estimating [cardinality](#) and other [statistics](#) for an indexed column, such as those calculated by [ANALYZE TABLE](#). Increasing the value improves the accuracy of index statistics, which can improve the [query execution plan](#), at the expense of increased I/O during the execution of [ANALYZE TABLE](#) for an InnoDB table. For more information, see [Section 17.8.10.1, “Configuring Persistent Optimizer Statistics Parameters”](#).

Note

Setting a high value for [innodb_stats_persistent_sample_pages](#) could result in lengthy [ANALYZE TABLE](#) execution time. To estimate the number of database pages accessed by [ANALYZE TABLE](#), see [Section 17.8.10.3, “Estimating ANALYZE TABLE Complexity for InnoDB Tables”](#).

[innodb_stats_persistent_sample_pages](#) only applies when [innodb_stats_persistent](#) is enabled for a table; when [innodb_stats_persistent](#) is disabled, [innodb_stats_transient_sample_pages](#) applies instead.

[innodb_stats_transient_sample_pages](#)

Command-Line Format	--innodb-stats-transient-sample-pages=#
System Variable	innodb_stats_transient_sample_pages
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	8
Minimum Value	1
Maximum Value	18446744073709551615

The number of index [pages](#) to sample when estimating [cardinality](#) and other [statistics](#) for an indexed column, such as those calculated by [ANALYZE TABLE](#). The default value is 8. Increasing the value improves the accuracy of index statistics, which can improve the [query execution plan](#), at the expense of increased I/O when opening an InnoDB table or recalculating statistics. For more information, see [Section 17.8.10.2, “Configuring Non-Persistent Optimizer Statistics Parameters”](#).

Note

Setting a high value for [innodb_stats_transient_sample_pages](#) could result in lengthy [ANALYZE TABLE](#) execution time. To estimate the number of database pages accessed by [ANALYZE TABLE](#), see [Section 17.8.10.3, “Estimating ANALYZE TABLE Complexity for InnoDB Tables”](#).

[innodb_stats_transient_sample_pages](#) only applies when [innodb_stats_persistent](#) is disabled for a table; when [innodb_stats_persistent](#) is enabled, [innodb_stats_persistent_sample_pages](#) applies instead. Takes the place of [innodb_stats_sample_pages](#) that was removed in MySQL 8.0. For more information, see [Section 17.8.10.2, “Configuring Non-Persistent Optimizer Statistics Parameters”](#).

[innodb_status_output](#)

Command-Line Format	--innodb-status-output[={OFF ON}]
System Variable	innodb_status_output
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	OFF

Enables or disables periodic output for the standard InnoDB Monitor. Also used in combination with [innodb_status_output_locks](#) to enable or disable periodic output for the InnoDB Lock Monitor. For more information, see [Section 17.17.2, “Enabling InnoDB Monitors”](#).

[innodb_status_output_locks](#)

Command-Line Format	--innodb-status-output-locks[={OFF ON}]
System Variable	innodb_status_output_locks
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	OFF

Enables or disables the InnoDB Lock Monitor. When enabled, the InnoDB Lock Monitor prints additional information about locks in `SHOW ENGINE INNODB STATUS` output and in periodic output printed to the MySQL error log. Periodic output for the InnoDB Lock Monitor is printed as part of the standard InnoDB Monitor output. The standard InnoDB Monitor must therefore be enabled for the InnoDB Lock Monitor to print data to the MySQL error log periodically. For more information, see [Section 17.17.2, “Enabling InnoDB Monitors”](#).

[innodb_strict_mode](#)

Command-Line Format	--innodb-strict-mode[={OFF ON}]
System Variable	innodb_strict_mode
Scope	Global, Session
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	ON

When [innodb_strict_mode](#) is enabled, InnoDB returns errors rather than warnings when checking for invalid or incompatible table options.

It checks that KEY_BLOCK_SIZE, ROW_FORMAT, DATA DIRECTORY, TEMPORARY, and TABLESPACE options are compatible with each other and other settings.

innodb_strict_mode=ON also enables a row size check when creating or altering a table, to prevent INSERT or UPDATE from failing due to the record being too large for the selected page size.

You can enable or disable [innodb_strict_mode](#) on the command line when starting mysqld, or in a MySQL [configuration file](#). You can also enable or disable [innodb_strict_mode](#) at runtime with the statement SET [GLOBAL|SESSION] innodb_strict_mode=mode, where mode is either ON or OFF. Changing the GLOBAL setting requires privileges sufficient to set global system variables (see [Section 7.1.9.1, "System Variable Privileges"](#)) and affects the operation of all clients that subsequently connect. Any client can change the SESSION setting for [innodb_strict_mode](#), and the setting affects only that client.

Setting the session value of this system variable is a restricted operation. The session user must have privileges sufficient to set restricted session variables. See [Section 7.1.9.1, "System Variable Privileges"](#).

[innodb_sync_array_size](#)

Command-Line Format	--innodb-sync-array-size=#
System Variable	innodb_sync_array_size
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	Integer
Default Value	1
Minimum Value	1
Maximum Value	1024

Defines the size of the mutex/lock wait array. Increasing the value splits the internal data structure used to coordinate threads, for higher concurrency in workloads with large numbers of waiting threads. This setting must be configured when the MySQL instance is starting up, and cannot be changed afterward. Increasing the value is recommended for workloads that frequently produce a large number of waiting threads, typically greater than 768.

[innodb_sync_spin_loops](#)

Command-Line Format	--innodb-sync-spin-loops=#
System Variable	innodb_sync_spin_loops
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	30
Minimum Value	0
Maximum Value	4294967295

The number of times a thread waits for an InnoDB mutex to be freed before the thread is suspended.

[innodb_sync_debug](#)

Command-Line Format	--innodb-sync-debug[={OFF ON}]
System Variable	innodb_sync_debug
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	Boolean
Default Value	OFF

Enables sync debug checking for the InnoDB storage engine. This option is only available if debugging support is compiled in using the [WITH_DEBUG CMake](#) option.

[innodb_table_locks](#)

Command-Line Format	--innodb-table-locks[={OFF ON}]
----------------------------	---------------------------------

System Variable	innodb_table_locks
Scope	Global, Session
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	ON

If `autocommit = 0`, InnoDB honors [LOCK TABLES](#); MySQL does not return from `LOCK TABLES ... WRITE` until all other threads have released all their locks to the table. The default value of [innodb_table_locks](#) is 1, which means that [LOCK TABLES](#) causes InnoDB to lock a table internally if `autocommit = 0`.

[innodb_table_locks = 0](#) has no effect for tables locked explicitly with [LOCK TABLES ... WRITE](#). It does have an effect for tables locked for read or write by [LOCK TABLES ... WRITE](#) implicitly (for example, through triggers) or by [LOCK TABLES ... READ](#).

For related information, see [Section 17.7, “InnoDB Locking and Transaction Model”](#).

[innodb_temp_data_file_path](#)

Command-Line Format	<code>--innodb-temp-data-file-path=file_name</code>
System Variable	innodb_temp_data_file_path
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	String
Default Value	<code>ibtmp1:12M:autoextend</code>

Defines the relative path, name, size, and attributes of global temporary tablespace data files. The global temporary tablespace stores rollback segments for changes made to user-created temporary tables.

If no value is specified for [innodb_temp_data_file_path](#), the default behavior is to create a single auto-extending data file named `ibtmp1` in the [innodb_data_home_dir](#) directory. The initial file size is slightly larger than 12MB.

The syntax for a global temporary tablespace data file specification includes the file name, file size, and `autoextend` and `max` attributes:

```
file_name:file_size[:autoextend[:max:max_file_size]]
```

The global temporary tablespace data file cannot have the same name as another InnoDB data file. Any inability or error creating the global temporary tablespace data file is treated as fatal and server startup is refused.

File sizes are specified in KB, MB, or GB by appending `K`, `M` or `G` to the size value. The sum of file sizes must be slightly larger than 12MB.

The size limit of individual files is determined by the operating system. File size can be more than 4GB on operating systems that support large files. Use of raw disk partitions for global temporary tablespace data files is not supported.

The `autoextend` and `max` attributes can be used only for the data file specified last in the [innodb_temp_data_file_path](#) setting. For example:

```
[mysqld]
innodb_temp_data_file_path=ibtmp1:50M;ibtmp2:12M:autoextend:max:500M
```

The `autoextend` option causes the data file to automatically increase in size when it runs out of free space. The `autoextend` increment is 64MB by default. To modify the increment, change the [innodb_autoextend_increment](#) variable setting.

The directory path for global temporary tablespace data files is formed by concatenating the paths defined by [innodb_data_home_dir](#) and [innodb_temp_data_file_path](#).

Before running InnoDB in read-only mode, set [innodb_temp_data_file_path](#) to a location outside of the data directory. The path must be relative to the data directory. For example:

```
--innodb-temp-data-file-path=../../tmp/ibtmp1:12M:autoextend
```

For more information, see [Global Temporary Tablespace](#).

[innodb_temp_tablespaces_dir](#)

Command-Line Format	<code>--innodb-temp-tablespaces-dir=dir_name</code>
System Variable	innodb_temp_tablespaces_dir
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	Directory name
Default Value	<code>#innodb_temp</code>

Defines the location where InnoDB creates a pool of session temporary tablespaces at startup. The default location is the #innodb_temp directory in the data directory. A fully qualified path or path relative to the data directory is permitted.

Session temporary tablespaces always store user-created temporary tables and internal temporary tables created by the optimizer using InnoDB. (Previously, the on-disk storage engine for internal temporary tables was determined by the [internal tmp disk storage engine](#) system variable, which is no longer supported. See [Storage Engine for On-Disk Internal Temporary Tables](#).)

For more information, see [Session Temporary Tablespaces](#).

[innodb_thread_concurrency](#)

Command-Line Format	--innodb-thread-concurrency=#
System Variable	innodb_thread_concurrency
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	0
Minimum Value	0
Maximum Value	1000

Defines the maximum number of threads permitted inside of InnoDB. A value of 0 (the default) is interpreted as infinite concurrency (no limit). This variable is intended for performance tuning on high concurrency systems.

InnoDB tries to keep the number of threads inside InnoDB less than or equal to the [innodb_thread_concurrency](#) limit. Threads waiting for locks are not counted in the number of concurrently executing threads.

The correct setting depends on workload and computing environment. Consider setting this variable if your MySQL instance shares CPU resources with other applications or if your workload or number of concurrent users is growing. Test a range of values to determine the setting that provides the best performance. [innodb_thread_concurrency](#) is a dynamic variable, which permits experimenting with different settings on a live test system. If a particular setting performs poorly, you can quickly set [innodb_thread_concurrency](#) back to 0.

Use the following guidelines to help find and maintain an appropriate setting:

If the number of concurrent user threads for a workload is consistently small and does not affect performance, set [innodb_thread_concurrency=0](#) (no limit).

If your workload is consistently heavy or occasionally spikes, set an [innodb_thread_concurrency](#) value and adjust it until you find the number of threads that provides the best performance. For example, suppose that your system typically has 40 to 50 users, but periodically the number increases to 60, 70, or more. Through testing, you find that performance remains largely stable with a limit of 80 concurrent users. In this case, set [innodb_thread_concurrency](#) to 80.

If you do not want InnoDB to use more than a certain number of virtual CPUs for user threads (20 virtual CPUs, for example), set [innodb_thread_concurrency](#) to this number (or possibly lower, depending on performance testing). If your goal is to isolate MySQL from other applications, consider binding the mysqld process exclusively to the virtual CPUs. Be aware, however, that exclusive binding can result in non-optimal hardware usage if the mysqld process is not consistently busy. In this case, you can bind the mysqld process to the virtual CPUs but allow other applications to use some or all of the virtual CPUs.

Note

From an operating system perspective, using a resource management solution to manage how CPU time is shared among applications may be preferable to binding the mysqld process. For example, you could assign 90% of virtual CPU time to a given application while other critical processes *are not* running, and scale that value back to 40% when other critical processes *are* running.

In some cases, the optimal [innodb_thread_concurrency](#) setting can be smaller than the number of virtual CPUs.

An [innodb_thread_concurrency](#) value that is too high can cause performance regression due to increased contention on system internals and resources.

Monitor and analyze your system regularly. Changes to workload, number of users, or computing environment may require that you adjust the [innodb_thread_concurrency](#) setting.

A value of 0 disables the queries inside InnoDB and queries in queue counters in the ROW OPERATIONS section of SHOW ENGINE INNODB STATUS output.

For related information, see [Section 17.8.4, “Configuring Thread Concurrency for InnoDB”](#).

[innodb_thread_sleep_delay](#)

Command-Line Format	--innodb-thread-sleep-delay=#
System Variable	innodb_thread_sleep_delay
Scope	Global

Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	10000
Minimum Value	0
Maximum Value	1000000
Unit	microseconds

How long InnoDB threads sleep before joining the InnoDB queue, in microseconds. The default value is 10000. A value of 0 disables sleep. You can set [innodb_adaptive_max_sleep_delay](#) to the highest value you would allow for [innodb_thread_sleep_delay](#), and InnoDB automatically adjusts [innodb_thread_sleep_delay](#) up or down depending on current thread-scheduling activity. This dynamic adjustment helps the thread scheduling mechanism to work smoothly during times when the system is lightly loaded or when it is operating near full capacity.

For more information, see [Section 17.8.4, “Configuring Thread Concurrency for InnoDB”](#).

[innodb_tmpdir](#)

Command-Line Format	<code>--innodb-tmpdir=dir_name</code>
System Variable	innodb_tmpdir
Scope	Global, Session
Dynamic	Yes
SET VAR Hint Applies	No
Type	Directory name
Default Value	NULL

Used to define an alternate directory for temporary sort files created during online [ALTER TABLE](#) operations that rebuild the table.

Online [ALTER TABLE](#) operations that rebuild the table also create an *intermediate* table file in the same directory as the original table. The [innodb_tmpdir](#) option is not applicable to intermediate table files.

A valid value is any directory path other than the MySQL data directory path. If the value is NULL (the default), temporary files are created MySQL temporary directory (\$TMPDIR on Unix, %TEMP% on Windows, or the directory specified by the `--tmpdir` configuration option). If a directory is specified, existence of the directory and permissions are only checked when [innodb_tmpdir](#) is configured using a [SET](#) statement. If a symlink is provided in a directory string, the symlink is resolved and stored as an absolute path. The path should not exceed 512 bytes. An online [ALTER TABLE](#) operation reports an error if [innodb_tmpdir](#) is set to an invalid directory. [innodb_tmpdir](#) overrides the MySQL [tmpdir](#) setting but only for online [ALTER TABLE](#) operations.

The FILE privilege is required to configure [innodb_tmpdir](#).

The [innodb_tmpdir](#) option was introduced to help avoid overflowing a temporary file directory located on a `tmpfs` file system. Such overflows could occur as a result of large temporary sort files created during online [ALTER TABLE](#) operations that rebuild the table.

In replication environments, only consider replicating the [innodb_tmpdir](#) setting if all servers have the same operating system environment. Otherwise, replicating the [innodb_tmpdir](#) setting could result in a replication failure when running online [ALTER TABLE](#) operations that rebuild the table. If server operating environments differ, it is recommended that you configure [innodb_tmpdir](#) on each server individually.

For more information, see [Section 17.12.3, “Online DDL Space Requirements”](#). For information about online [ALTER TABLE](#) operations, see [Section 17.12, “InnoDB and Online DDL”](#).

[innodb_trx_purge_view_update_only_debug](#)

Command-Line Format	<code>--innodb-trx-purge-view-update-only-debug[={OFF ON}]</code>
System Variable	innodb_trx_purge_view_update_only_debug
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	OFF

Pauses purging of delete-marked records while allowing the purge view to be updated. This option artificially creates a situation in which the purge view is updated but purges have not yet been performed. This option is only available if debugging support is compiled in using the [WITH_DEBUG CMake](#) option.

[innodb_trx_rseg_n_slots_debug](#)

Command-Line Format	<code>--innodb-trx-rseg-n-slots-debug=#</code>
System Variable	innodb_trx_rseg_n_slots_debug
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	0
Minimum Value	0
Maximum Value	1024

Sets a debug flag that limits TRX_RSEG_N_SLOTS to a given value for the `trx_rsegf_undo_find_free` function that looks for free slots for undo log segments. This option is only available if debugging support is compiled in using the [WITH_DEBUG CMake](#) option.

[innodb_undo_directory](#)

Command-Line Format	<code>--innodb-undo-directory=dir_name</code>
System Variable	innodb_undo_directory
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	Directory name

The path where InnoDB creates undo tablespaces. Typically used to place undo tablespaces on a different storage device.

There is no default value (it is NULL). If the [innodb_undo_directory](#) variable is undefined, undo tablespaces are created in the data directory.

The default undo tablespaces (`innodb_undo_001` and `innodb_undo_002`) created when the MySQL instance is initialized always reside in the directory defined by the [innodb_undo_directory](#) variable.

Undo tablespaces created using [CREATE UNDO TABLESPACE](#) syntax are created in the directory defined by the [innodb_undo_directory](#) variable if a different path is not specified.

For more information, see [Section 17.6.3.4, “Undo Tablespaces”](#).

[innodb_undo_log_encrypt](#)

Command-Line Format	<code>--innodb-undo-log-encrypt[={OFF ON}]</code>
System Variable	innodb_undo_log_encrypt
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	OFF

Controls encryption of undo log data for tables encrypted using the InnoDB [data-at-rest encryption feature](#). Only applies to undo logs that reside in separate [undo tablespaces](#). See [Section 17.6.3.4, “Undo Tablespaces”](#). Encryption is not supported for undo log data that resides in the system tablespace. For more information, see [Undo Log Encryption](#).

[innodb_undo_log_truncate](#)

Command-Line Format	<code>--innodb-undo-log-truncate[={OFF ON}]</code>
System Variable	innodb_undo_log_truncate
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	ON

When enabled, undo tablespaces that exceed the threshold value defined by [innodb_max_undo_log_size](#) are marked for truncation. Only undo tablespaces can be truncated. Truncating undo logs that reside in the system tablespace is not supported. For truncation to occur, there must be at least two undo tablespaces.

The [innodb_purge_rseg_truncate_frequency](#) variable can be used to expedite truncation of undo tablespaces.

For more information, see [Truncating Undo Tablespaces](#).

[innodb_undo_tablespaces](#)

Command-Line Format	<code>--innodb-undo-tablespaces=#</code>
Deprecated	Yes
System Variable	innodb_undo_tablespaces
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Integer
Default Value	2
Minimum Value	2
Maximum Value	127

Defines the number of [undo tablespaces](#) used by InnoDB. The default and minimum value is 2.

Note

The [innodb_undo_tablespaces](#) variable is deprecated; setting it has no effect. You should expect it to be removed in a future MySQL release.

For more information, see [Section 17.6.3.4, “Undo Tablespaces”](#).

[innodb_use_fdatasync](#)

Command-Line Format	<code>--innodb-use-fdatasync[={OFF ON}]</code>
System Variable	innodb_use_fdatasync
Scope	Global
Dynamic	Yes
SET VAR Hint Applies	No
Type	Boolean
Default Value	ON

On platforms that support `fdatasync()` system calls, having [innodb_use_fdatasync](#) enabled permits using `fdatasync()` instead of `fsync()` system calls for operating system flushes. An `fdatasync()` call does not flush changes to file metadata unless required for subsequent data retrieval, providing a potential performance benefit.

A subset of [innodb_flush_method](#) settings such as `fsync`, `O_DSYNC`, and `O_DIRECT` use `fsync()` system calls. The [innodb_use_fdatasync](#) variable is applicable when using those settings.

Before MySQL 8.4, this option was disabled by default.

[innodb_use_native_aio](#)

Command-Line Format	<code>--innodb-use-native-aio[={OFF ON}]</code>
System Variable	innodb_use_native_aio
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	Boolean
Default Value	ON

Specifies whether to use the [asynchronous I/O](#) subsystem. This variable cannot be changed while the server is running. Normally, you do not need to configure this option, because it is enabled by default.

This feature improves the scalability of heavily I/O-bound systems, which typically show many pending reads/writes in `SHOW ENGINE INNODB STATUS` output.

Running with a large number of InnoDB I/O threads, and especially running multiple such instances on the same server machine, can exceed capacity limits on Linux systems. In this case, you may receive the following error:

```
EAGAIN: The specified maxevents exceeds the user's limit of available events.
```

You can typically address this error by writing a higher limit to `/proc/sys/fs/aio-max-nr`.

However, if a problem with the asynchronous I/O subsystem in the OS prevents InnoDB from starting, you can start the server with [innodb_use_native_aio=0](#). This option may also be disabled automatically during startup if InnoDB detects a potential problem such as a combination of `tmpdir` location, `tmpfs` file system, and Linux kernel that does not support AIO on `tmpfs`.

For more information, see [Section 17.8.6, “Using Asynchronous I/O on Linux”](#).

[innodb_validate_tablespace_paths](#)

Command-Line Format	<code>--innodb-validate-tablespace-paths[={OFF ON}]</code>
System Variable	innodb_validate_tablespace_paths
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	Boolean
Default Value	ON

Controls tablespace file path validation. At startup, InnoDB validates the paths of known tablespace files against tablespace file paths stored in the data dictionary in case tablespace files have been moved to a different location. The [innodb_validate_tablespace_paths](#) variable permits disabling tablespace path validation. This feature is intended for environments where tablespaces files are not moved. Disabling path validation improves startup time on systems with a large number of tablespace files.

Warning

Starting the server with tablespace path validation disabled after moving tablespace files can lead to undefined behavior.

For more information, see [Section 17.6.3.7, “Disabling Tablespace Path Validation”](#).

[innodb version](#)

The InnoDB version number. This is a legacy variable, the value is the same as the MySQL server [version](#).

[innodb write io threads](#)

Command-Line Format	--innodb-write-io-threads=#
System Variable	innodb write io threads
Scope	Global
Dynamic	No
SET VAR Hint Applies	No
Type	Integer
Default Value	4
Minimum Value	1
Maximum Value	64

The number of I/O threads for write operations in InnoDB. The default value is 4. Its counterpart for read threads is [innodb read io threads](#). For more information, see [Section 17.8.5, “Configuring the Number of Background InnoDB I/O Threads”](#). For general I/O tuning advice, see [Section 10.5.8, “Optimizing InnoDB Disk I/O”](#).

Note

On Linux systems, running multiple MySQL servers (typically more than 12) with default settings for [innodb read io threads](#), [innodb write io threads](#), and the Linux `aio-max-nr` setting can exceed system limits. Ideally, increase the `aio-max-nr` setting; as a workaround, you might reduce the settings for one or both of the MySQL variables.

Also take into consideration the value of [sync_binlog](#), which controls synchronization of the binary log to disk.

For general I/O tuning advice, see [Section 10.5.8, “Optimizing InnoDB Disk I/O”](#).

[PREV](#) [HOME](#) [UP](#) [NEXT](#)

[Related Documentation](#)

[MySQL 8.4 Release Notes](#)

[Download this Manual](#)

[PDF \(US Ltr\)](#) - 39.9Mb

[PDF \(A4\)](#) - 40.0Mb

[Man Pages \(TGZ\)](#) - 258.3Kb

[Man Pages \(Zip\)](#) - 365.2Kb

[Info \(Gzip\)](#) - 4.0Mb

[Info \(Zip\)](#) - 4.0Mb